

Side channel example: The Cold-Boot Attack

Lest We Remember

Cold-Boot Attacks Against Disk Encryption

J. Alex Halderman, Seth D. Schoen, Nadia Heninger,
William Clarkson, William Paul, Joseph A. Calandrino,
Ariel J. Feldman, Jacob Appelbaum, Edward W. Felten

Data Theft Threat



OS Access Control

Attacker's Computer

How Safe is a Stolen Laptop?



Unlocked
Unencrypted
No protection



Locked
Unencrypted
Minimal protection



Locked
Encrypted
Industry best practice

How Safe is a Stolen Laptop?

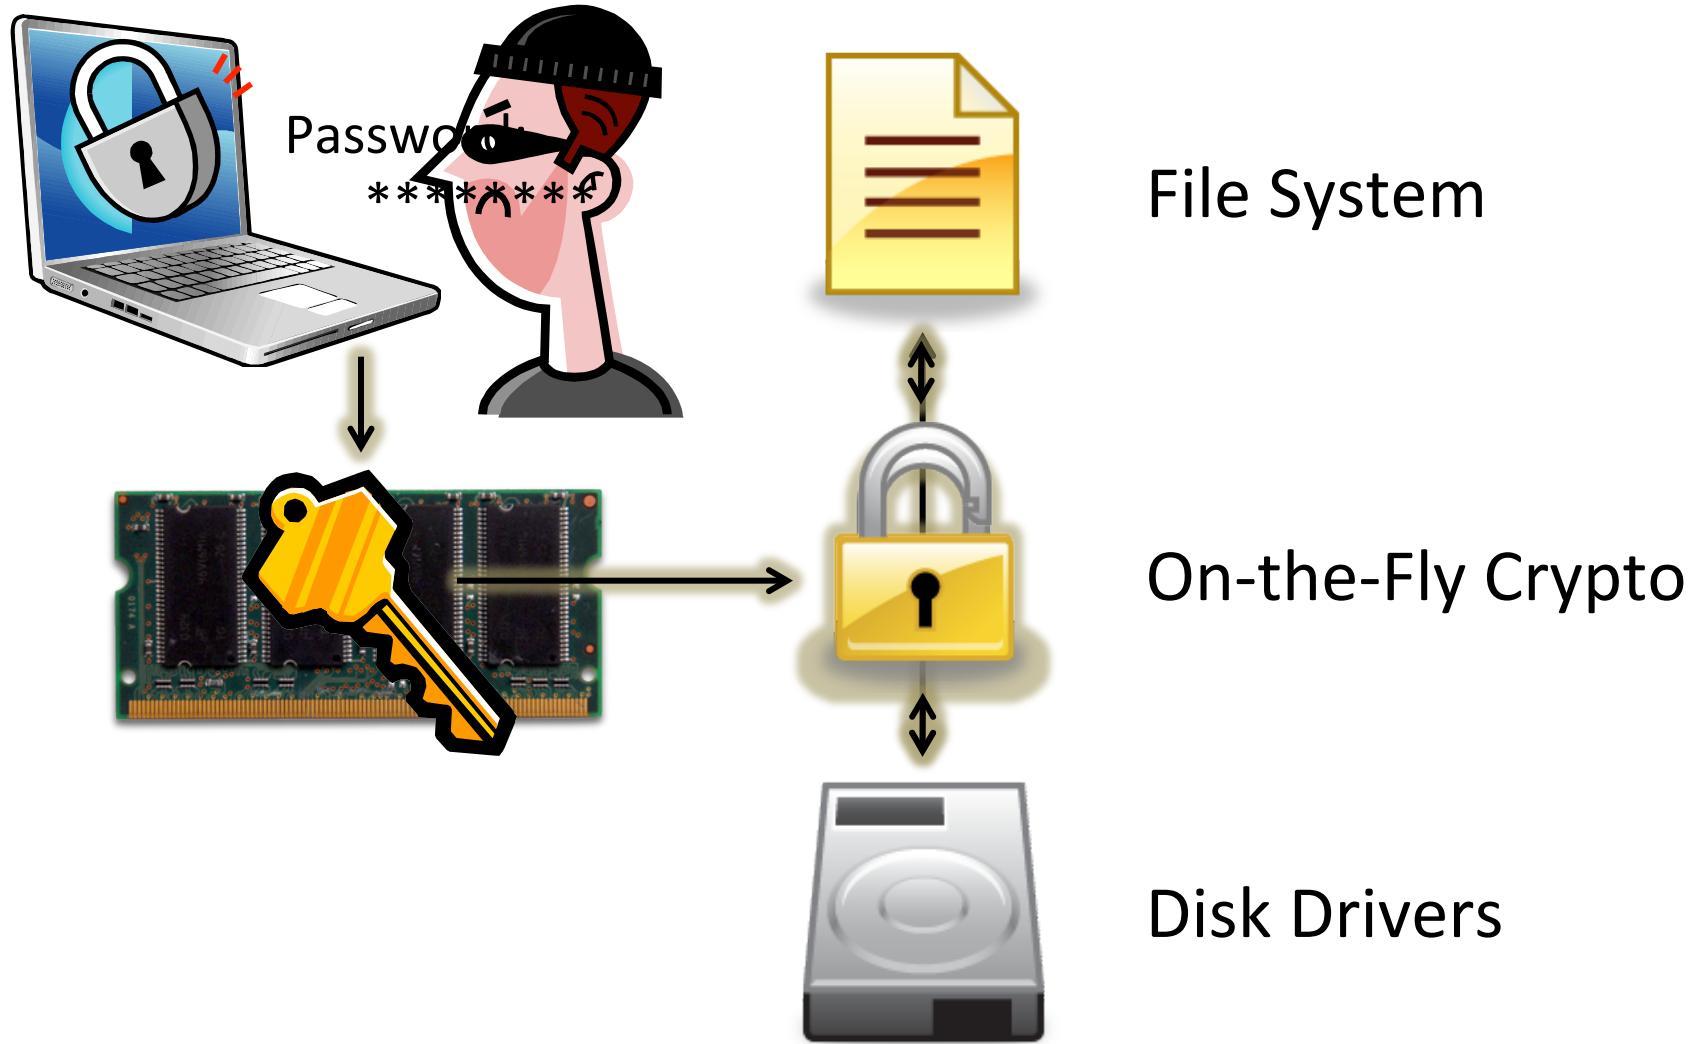
In 2007: **119,000** encrypted laptops with sensitive data lost at major US airports

Were they *really* protected?



**Locked
Encrypted**
Industry best practice

Disk Encryption Defense



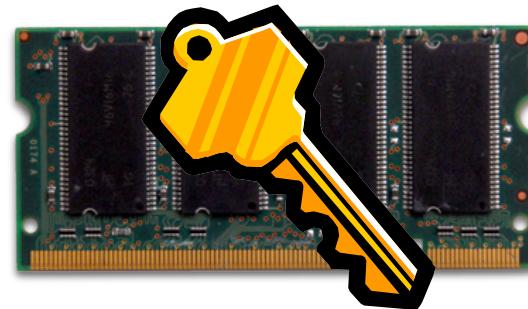
Common Attack Scenario



Security Assumptions

The encryption is strong

The OS protects the key in RAM

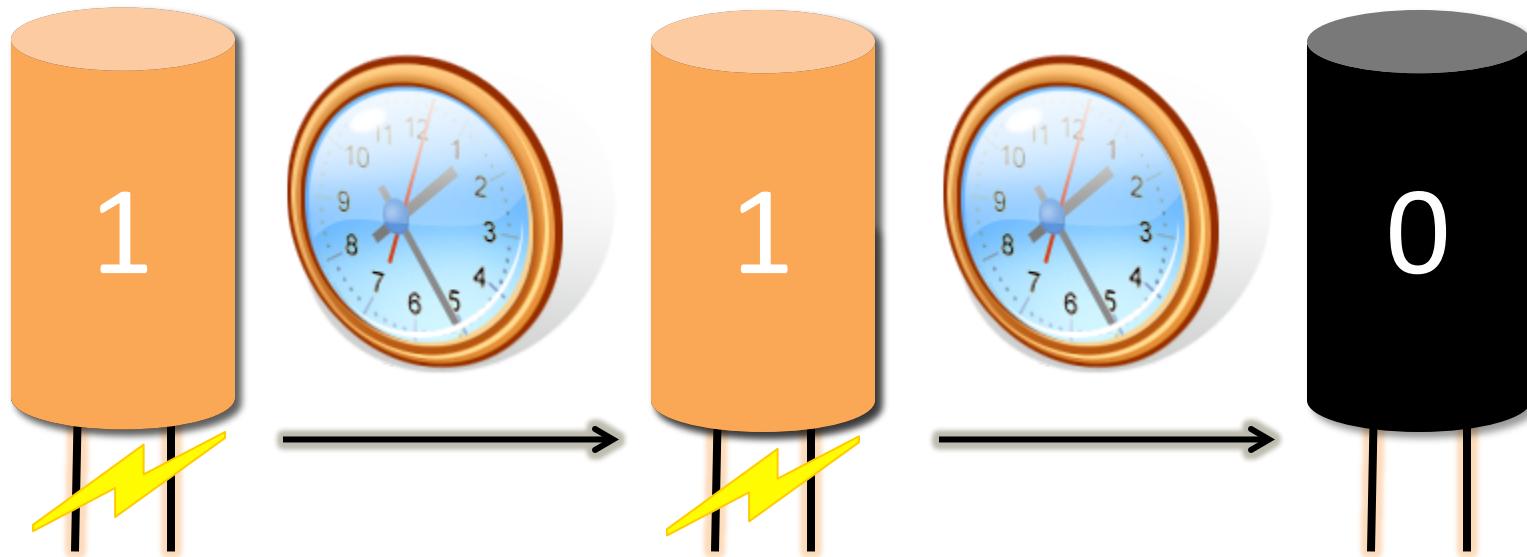


...the attacker might reboot to circumvent the OS, but since RAM is volatile, the key will be lost...

Right...?

Dynamic RAM Volatility

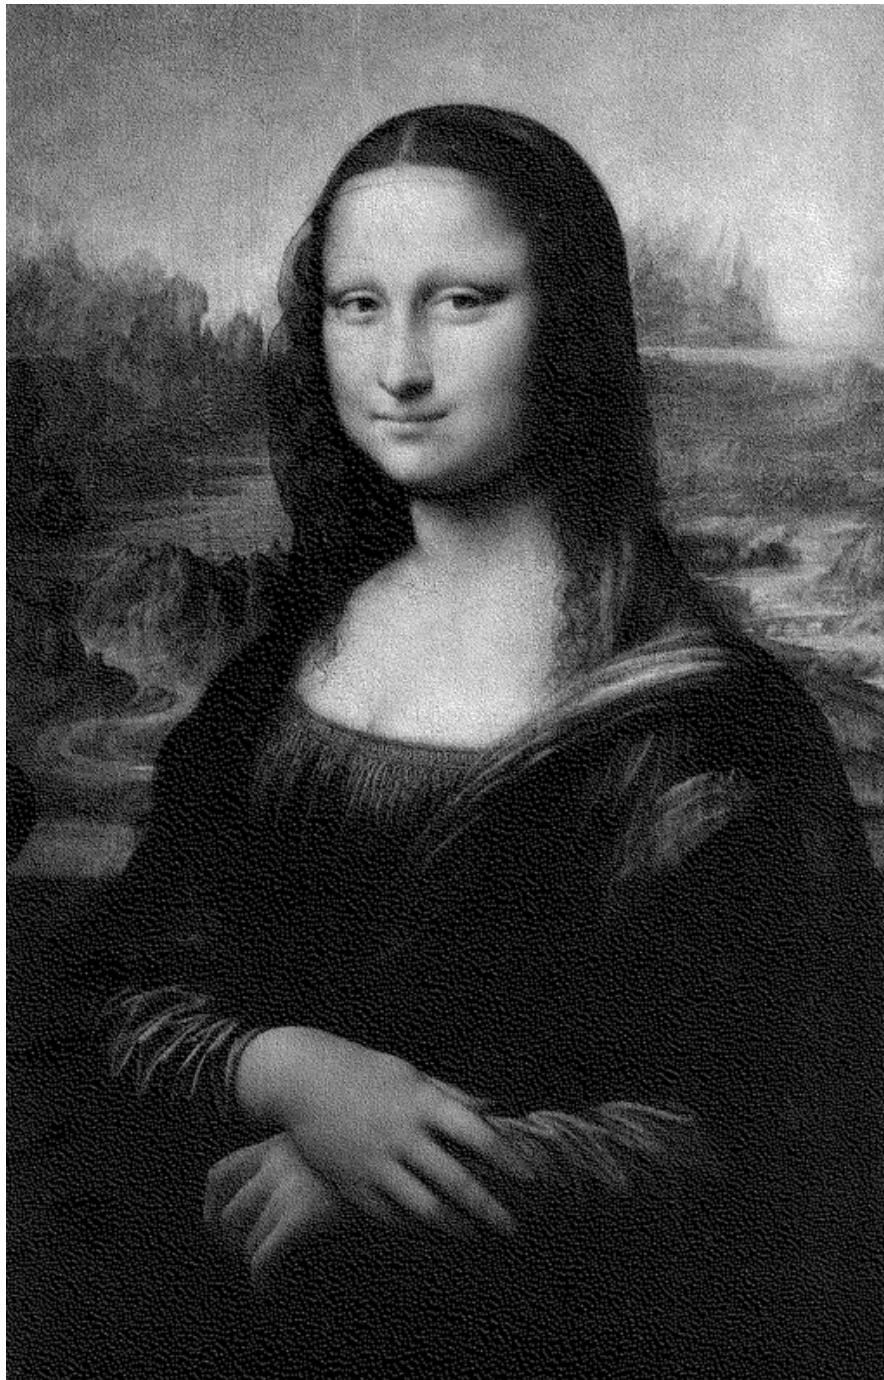
DRAM Cell
(Capacitor)



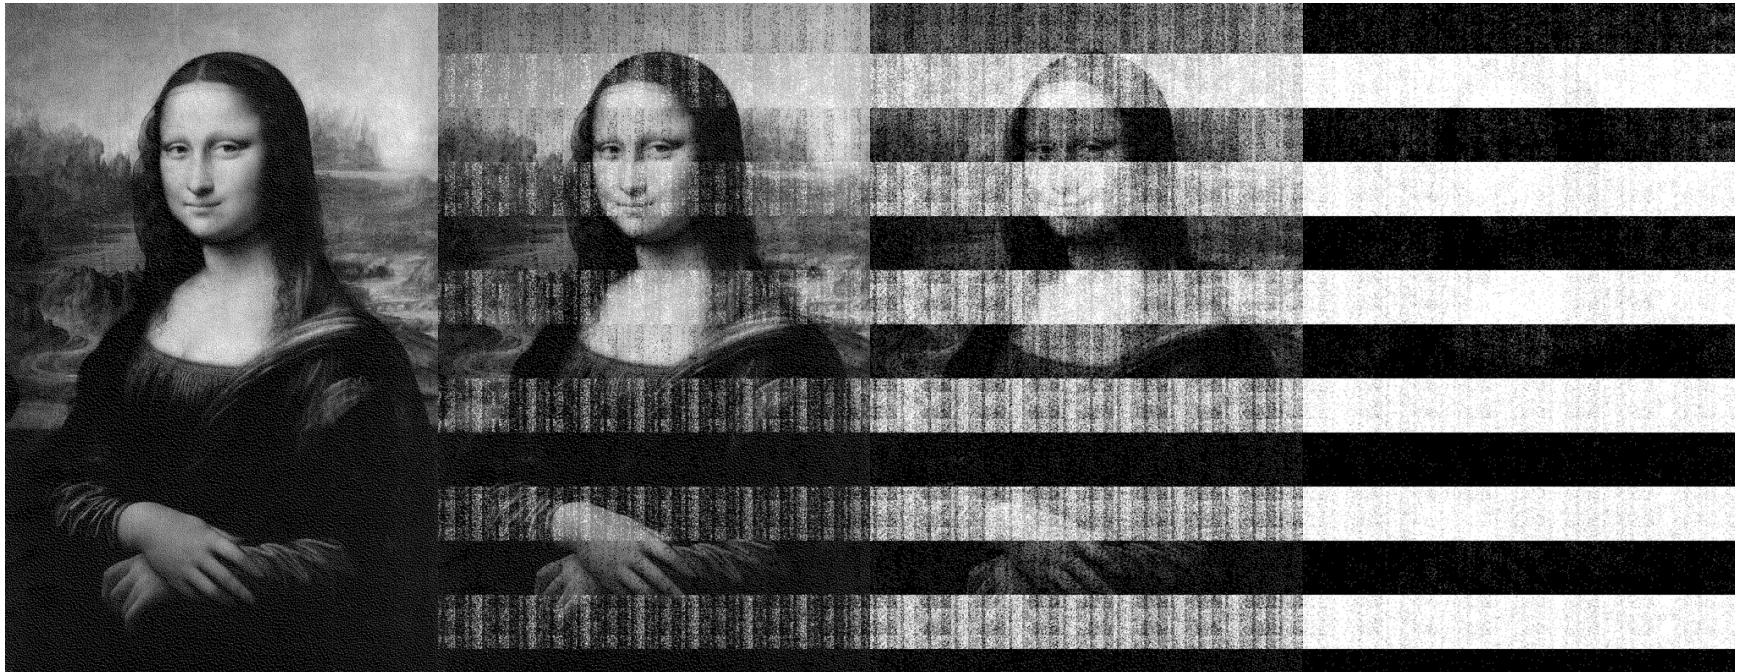
Write "1"
Security Assumption: Refresh
what if we don't refresh?
Data fades almost instantaneously without refresh
Any residual data is difficult to recover

Refresh Interval \approx 32 ms

(Read or Rewrite)



Decay After Cutting Power



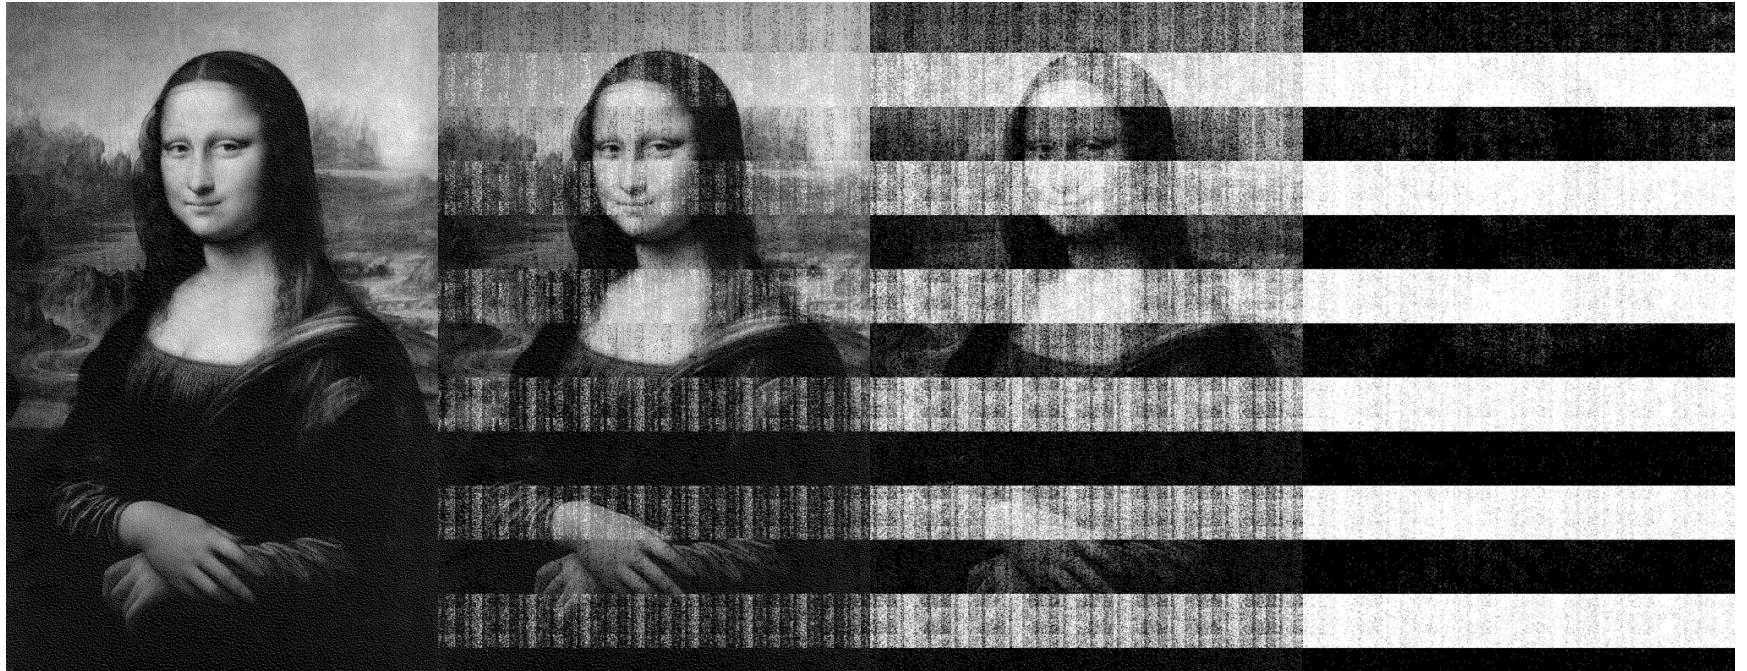
5 secs

30 secs

60 secs

300 secs

Decay After Cutting Power



5 secs

30 secs

60 secs

300 secs

~~DRAM data disappears almost instantaneously without refresh~~

Gradual

Unidirectional

Predictable

Capturing Residual Data

~~Any residual data is difficult to recover~~

Residual data can be captured easily, with no special equipment

Complication

Booting full OS overwrites large areas of RAM

Solution

Boot a small low-level program to write out (dump) memory content

Implementations

PXE Dump (9 KB)

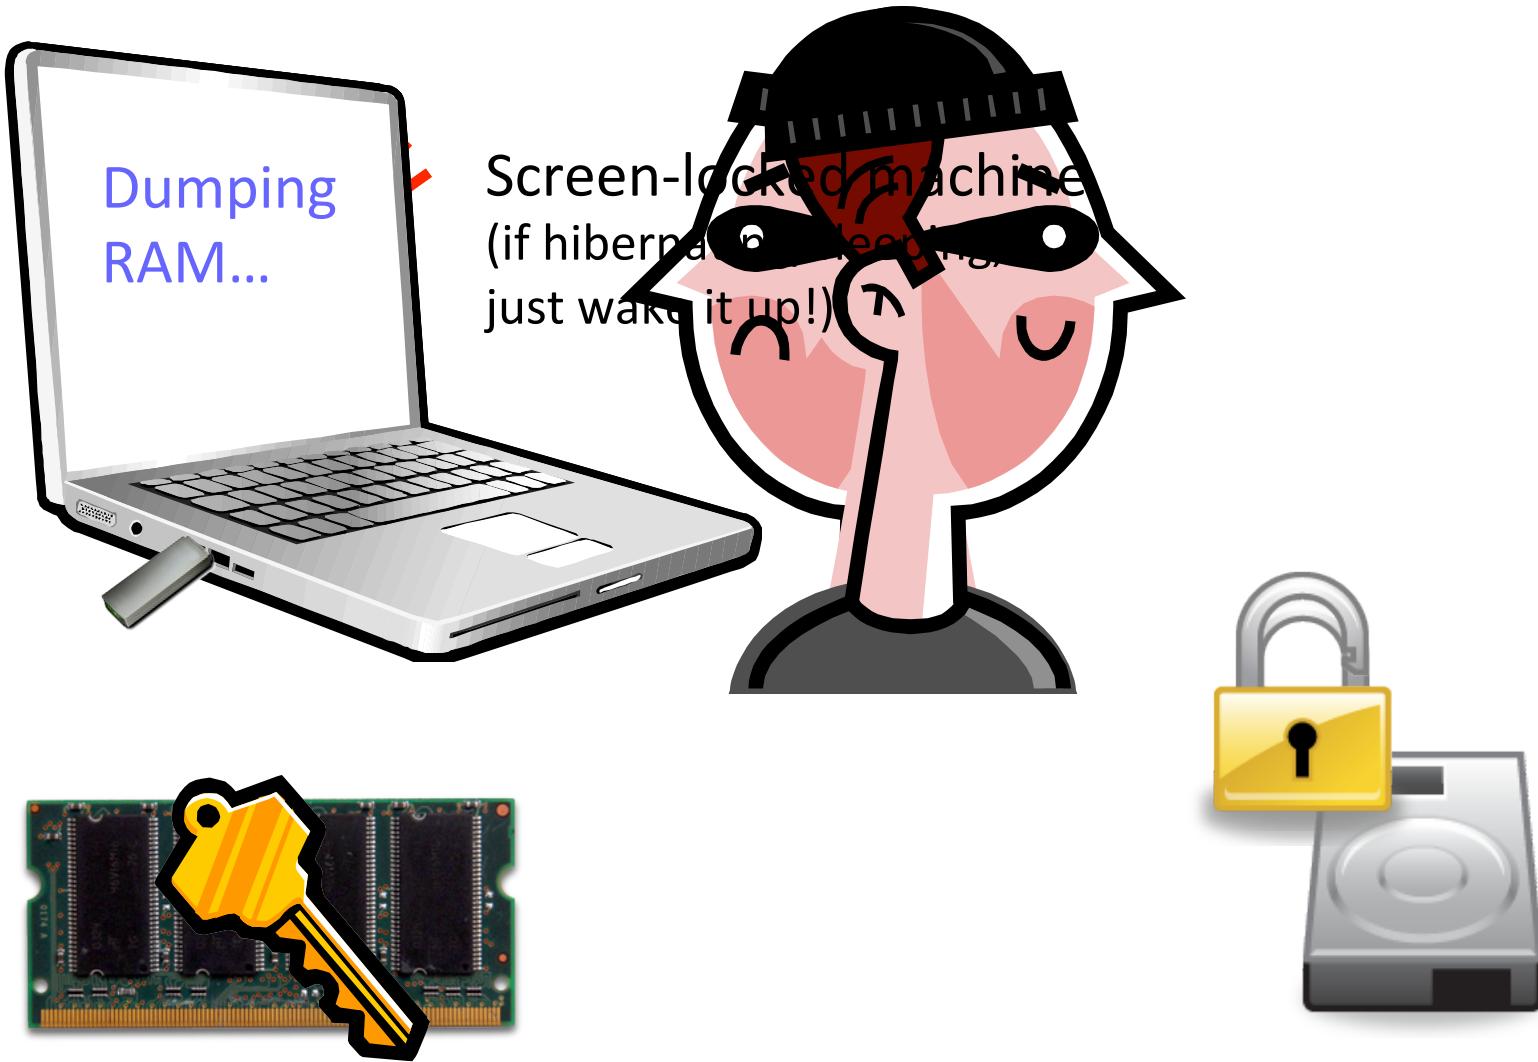
EFI Dump (10 KB)

USB Dump (22 KB)

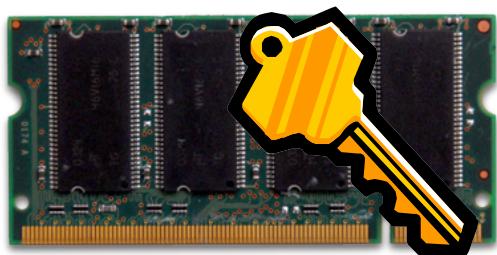
Delivering the Attack



Basic Cold-Boot Attack



What if the BIOS clears RAM?



Common in machines
that support ECC RAM

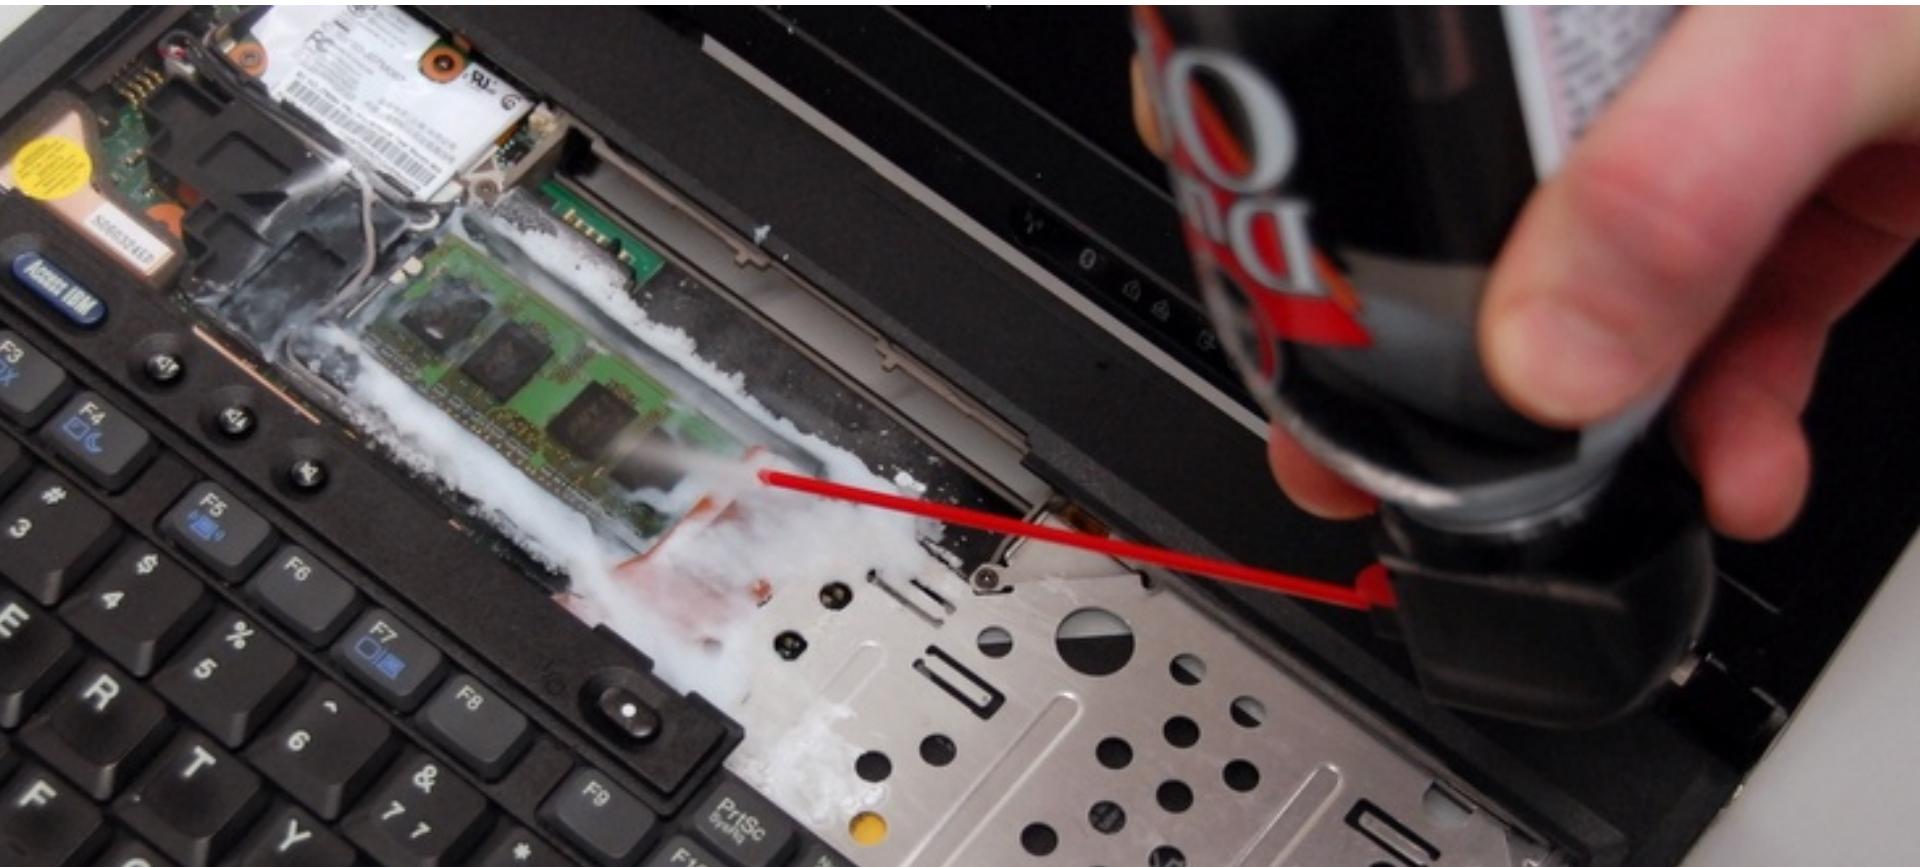
Advanced Cold-Boot Attack



Victim's Computer

Attacker's Computer

Slowing Decay By Cooling



-50°C

< 0.2% decay after 1 minute





SUN MON TUE WED THU FRI SAT
6:48:29
SEC

GOWalking
by SPORTLINE
WATER RESISTANT



Even Cooler

Liquid Nitrogen -196°C

< 0.17% decay after 1 hour

Not necessary in practice

Dealing with Bit Errors

Some bit errors inevitable, especially without cooling
(increasing with memory density)

Naïve Approach

Given corrupted K' , find K :

Brute-force search over low
Hamming distance to K'

e.g. 256-bit key with 10% error:

$> 2^{56}$ guesses (too slow!)

Insight

Most programs store
precomputed derivatives of K
(e.g. key schedules)

These derivatives contain
redundancy; we can treat
them as **error correcting codes**

Recovering Other Keys

AES, DES (key schedules)

LRW tweak keys (multiplication tables)

RSA private keys (primes P and Q)

Efficiency vs. Security

Finding Keys

Previous Methods

- Entropy methods
 - But, most high entropy data isn't key material
- Looking for data structure
 - But, need to know internal details about software

Our Method

- Focuses on key schedules
 - Precise and fully automatic, even with decay
 - Doesn't require analysis of target program

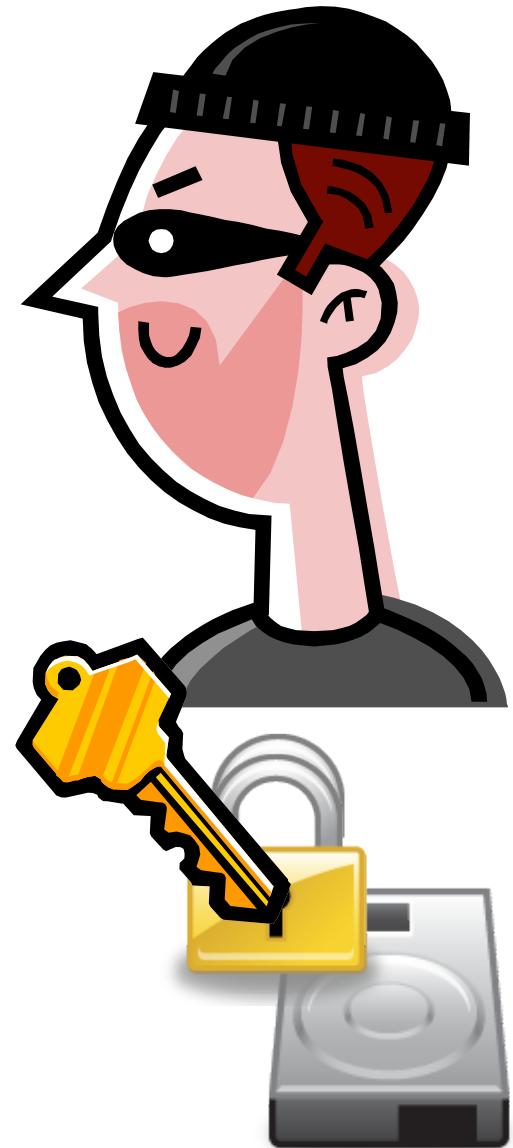
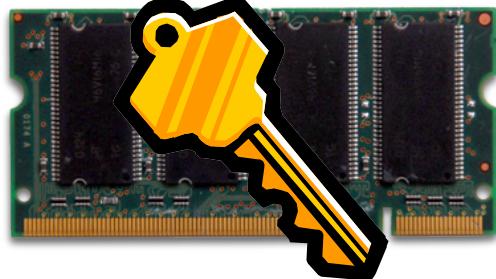
Finding AES Key Schedules

Iterate through each byte of memory –

1. Treat following region as an AES key schedule
2. For each word in the “schedule”:
 - Calculate correct value, assuming other bytes correct
 - Take Hamming distance to observed value
3. If total distance is low, output the key

Steps of Cold-Boot Attack

1. Extract memory
2. Locate keys in memory
3. Reconstruct decayed keys
4. Decrypt hard drive



Demonstrated Attacks

Windows BitLocker

Mac OS FileVault

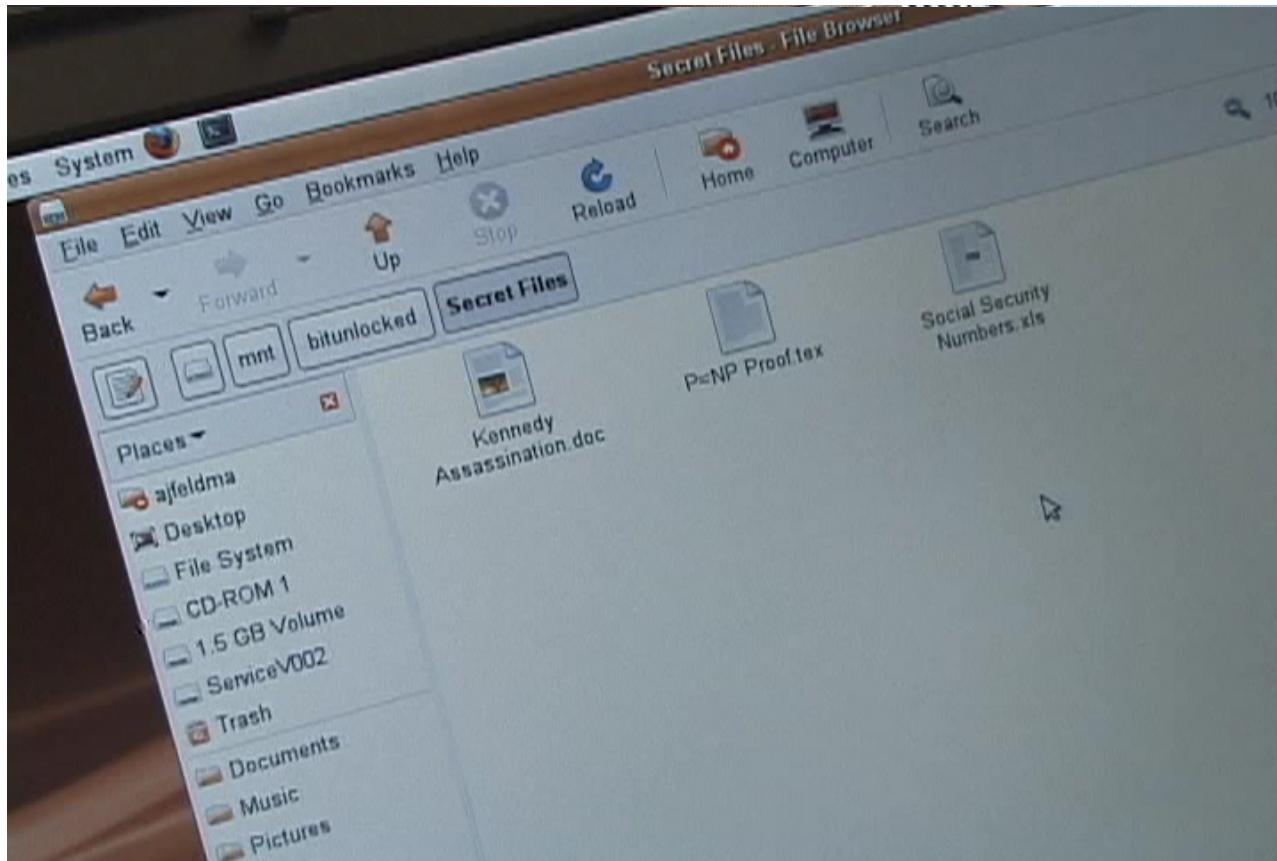
Linux dm-crypt

Linux LoopAES

TrueCrypt



“BitLocker, Meet BitUnlocker”



Demonstration of fully automated attack:
Connect USB drive, reboot, and browse files

Countermeasures

No Magic Bullet

Possible Mitigations

- Encrypt key in memory when screen-locked
- Avoid precomputation
- Fully encrypted memory
- Trusted Platform Module (TPM)

Broader Lessons

- Mistaken assumptions can compromise the security of otherwise good crypto
- Abstractions (e.g., hardware/software divide) can hide critical behaviors and requirements
- Ultimately, security is an empirical property of systems – experiments can support or overturn theory

**Covert channel example:
EC2**

Hey, You, Get Off of My Cloud: Exploring Information Leakage in Third-Party Compute Clouds, by

Ristenpart et al.

A simplified model of public cloud computing

Users run Virtual Machines (VMs) on cloud provider's infrastructure



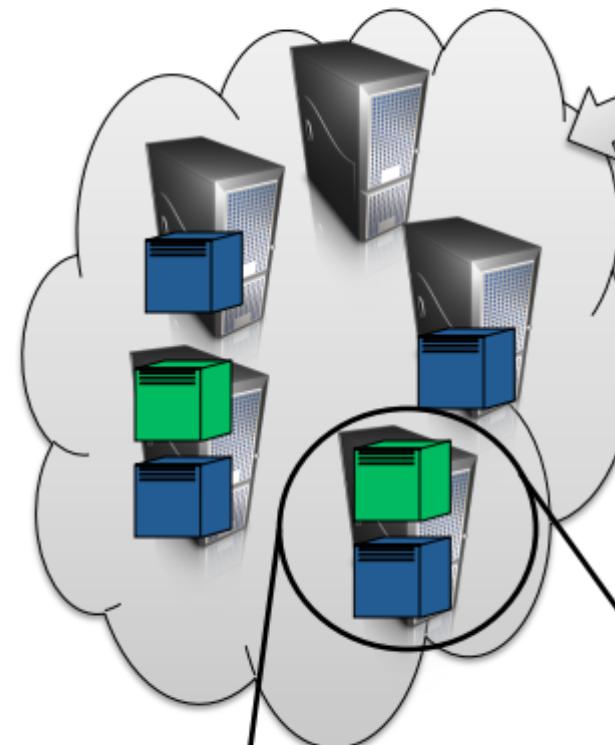
User A

virtual machines (VMs)

Three blue 3D cube icons representing virtual machines.

User B

virtual machines (VMs)

Two green 3D cube icons representing virtual machines.

Owned/operated
by cloud provider

Windows Azure

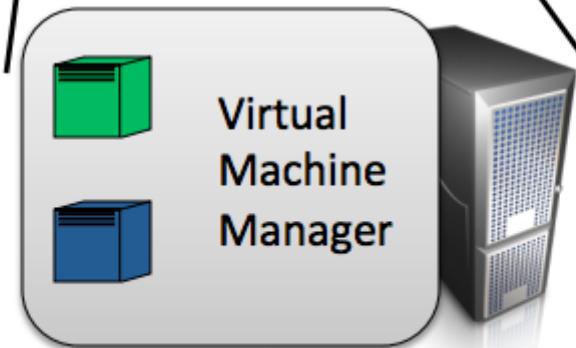
MOSSO™

amazon web services™

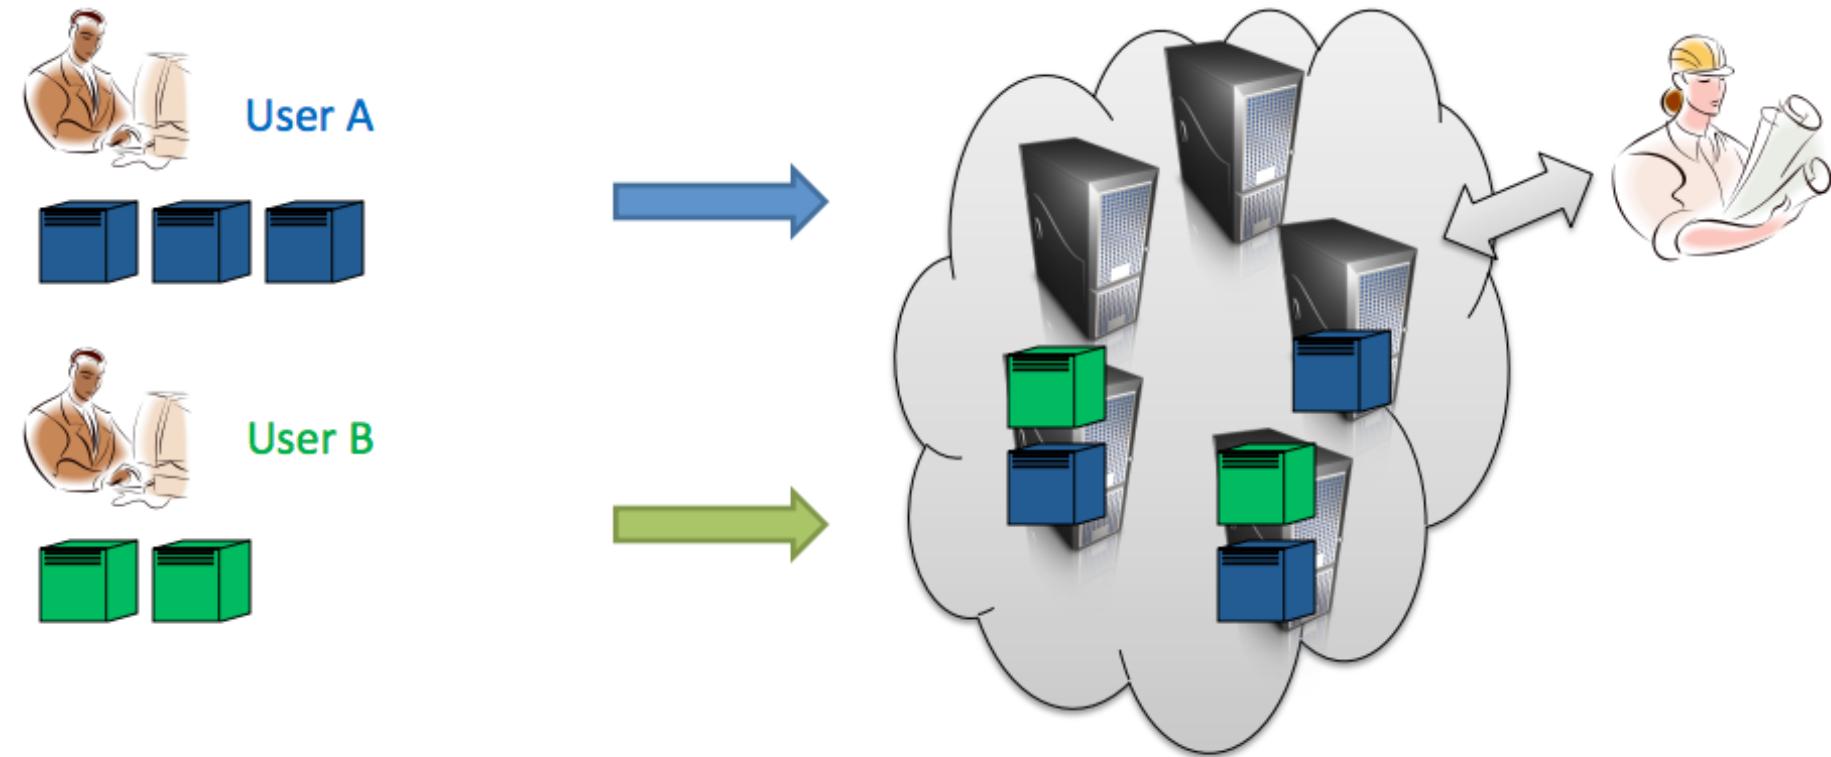
Multitenancy (users share physical resources)

Virtual Machine Manager (VMM)
manages physical server resources for VMs

To the VM should look like dedicated server



Trust models in public cloud computing



Users must trust third-party provider to

not spy on running VMs / data

secure infrastructure from external attackers

secure infrastructure from internal attackers

Outline of a more damaging approach:

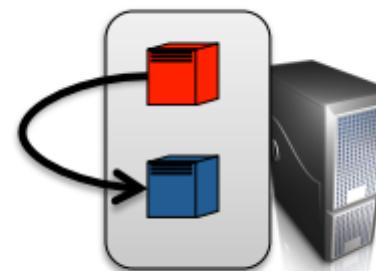
1) Cloud cartography

map internal infrastructure of cloud
map used to locate targets in cloud



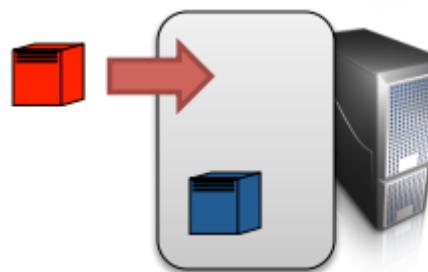
2) Checking for co-residence

check that VM is on same server as target
- network-based co-residence checks
- efficacy confirmed by covert channels



3) Achieving co-residence

brute forcing placement
instance flooding after target launches



Placement vulnerability:
attackers can knowingly achieve co-residence with target

4) Location-based attacks

side-channels, DoS, escape-from-VM



An Exploration of L2 Cache Covert Channels in Virtualized Environments

Yunjing Xu, Michael Bailey, Farnam Jahanian

University of Michigan

Kaustubh Joshi, Matti Hitunen, Rick Schlichting

AT&T Labs Research

Motivation

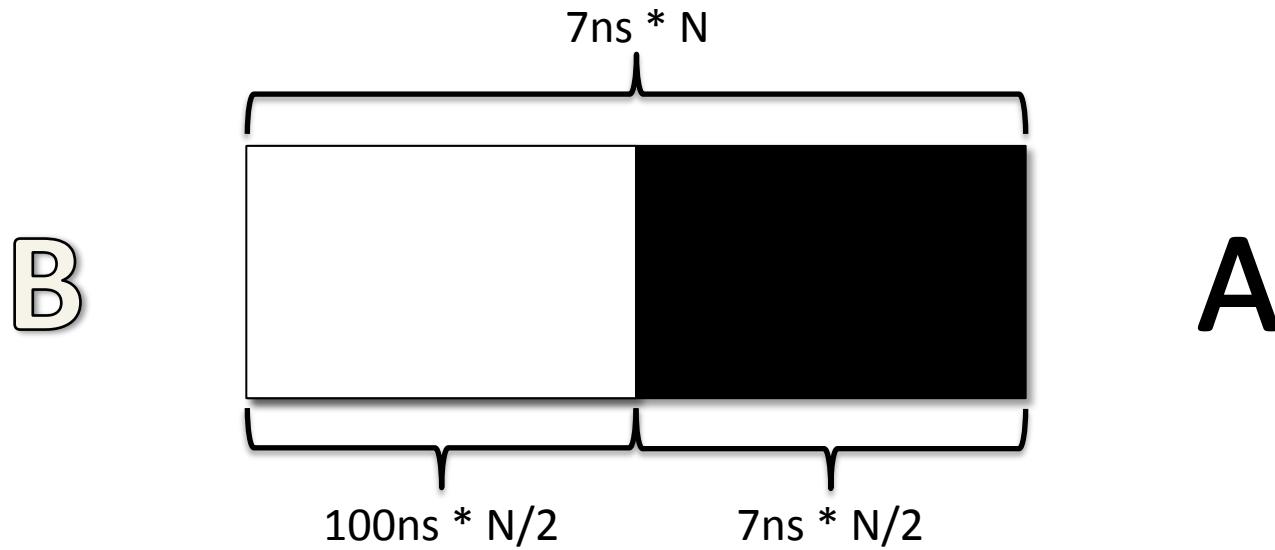
- The cloud environment (re)introduces shared hardware
- While powerful, this can be abused by attackers for information leakage
- Ristenpart et al. pioneering work in 2009
 - Forced VM co-location in EC2
 - Cross-VM information leakage
 - Covert channels
 - Side channels

Motivation

- Cross-VM covert channel using shared L2 cache
 - 0.2bps: proof-of-concept example
- What is the maximum achievable bit rate in practice over cross-VM covert channels using L2 cache?
- What factors influence the bit rate achievable in this scenario?

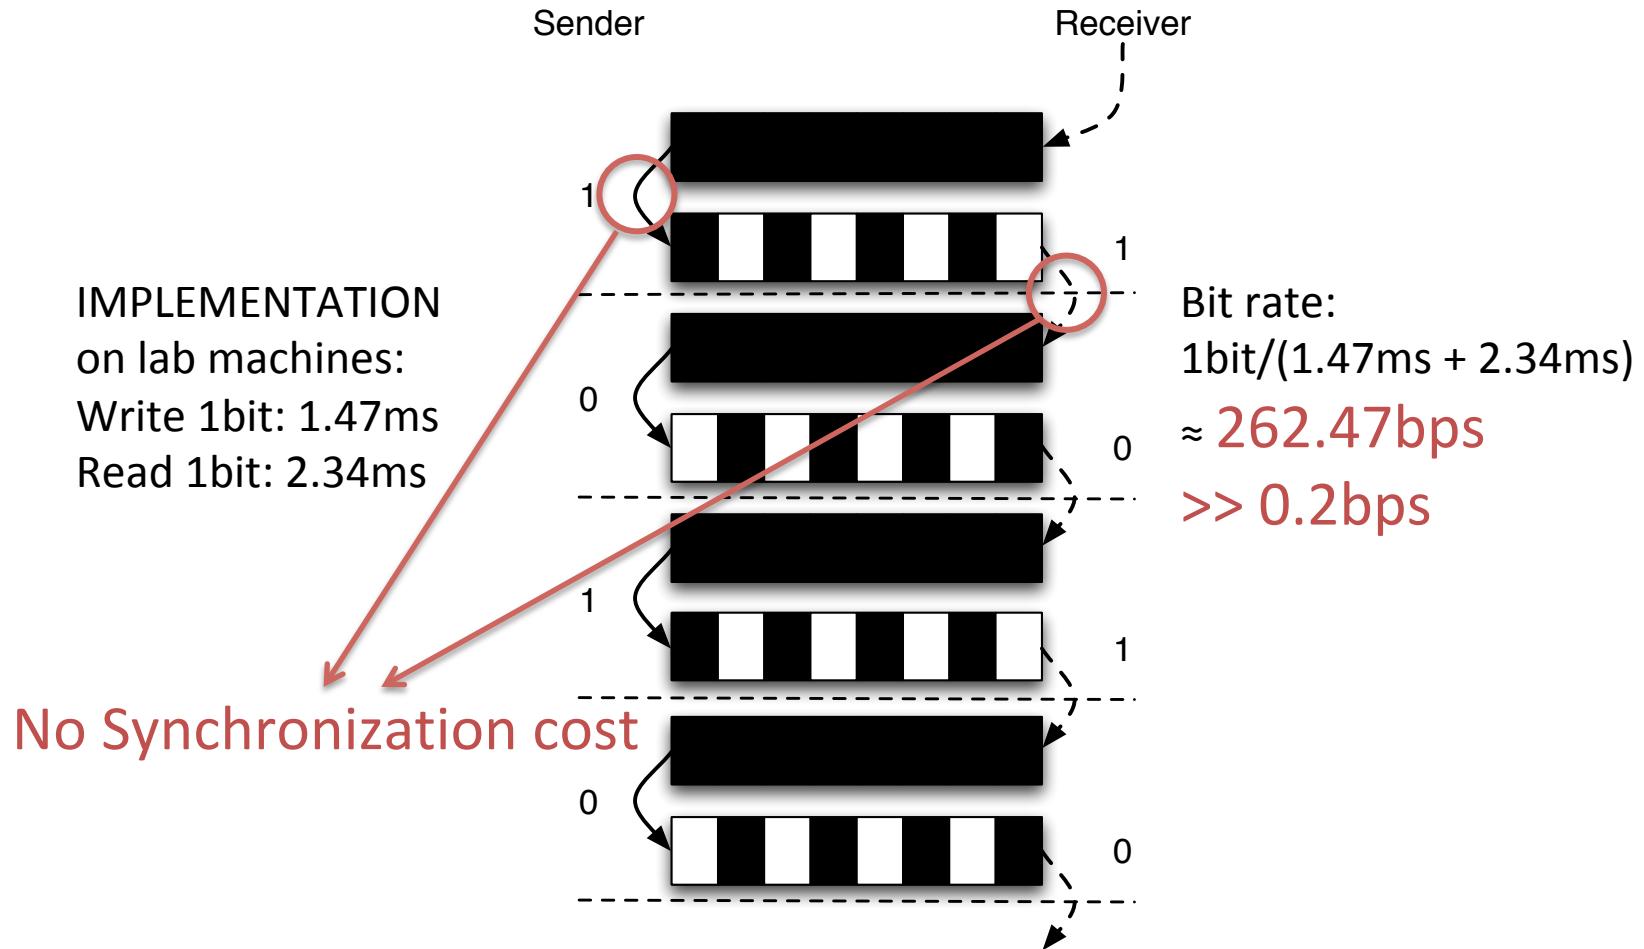
Cross-VM L2 Cache Covert Channel

- Processes on different VMs sharing a L2 cache can covertly communicate by manipulating the cache content



$$I_2 = \log 2 = 1\text{bit}$$
$$I_n = \log n$$

Bit Rate Estimation



Cost of Synchronization

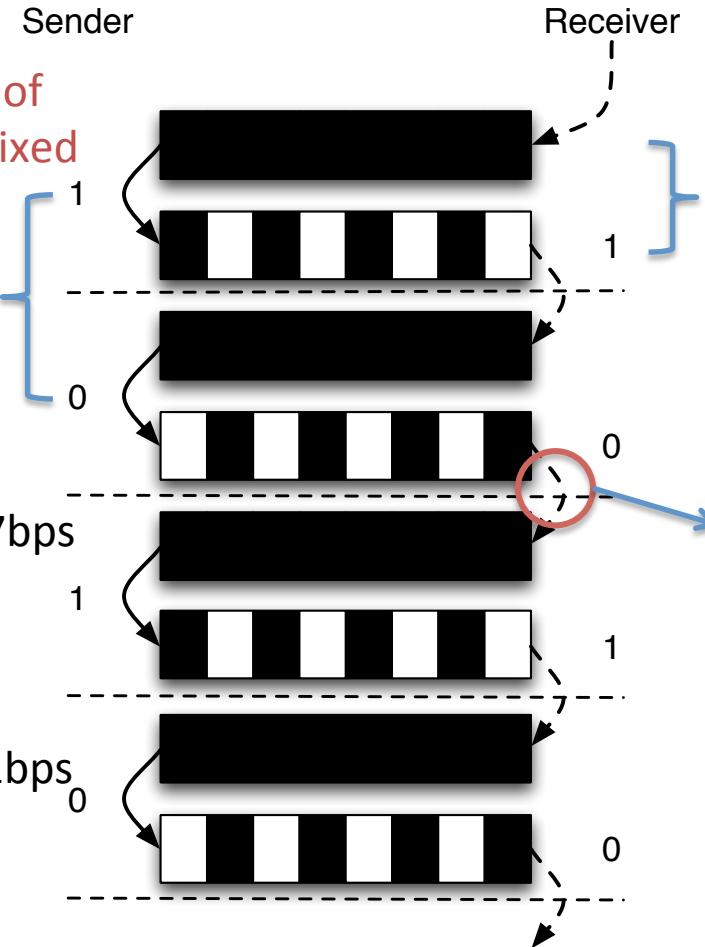
ASSUMPTION: The variation of reading time is small, if not fixed

Timing:
Sleep 3ms

$$1\text{bit}/(1.47 + 2.34)\text{ms} \approx 262.47\text{bps}$$



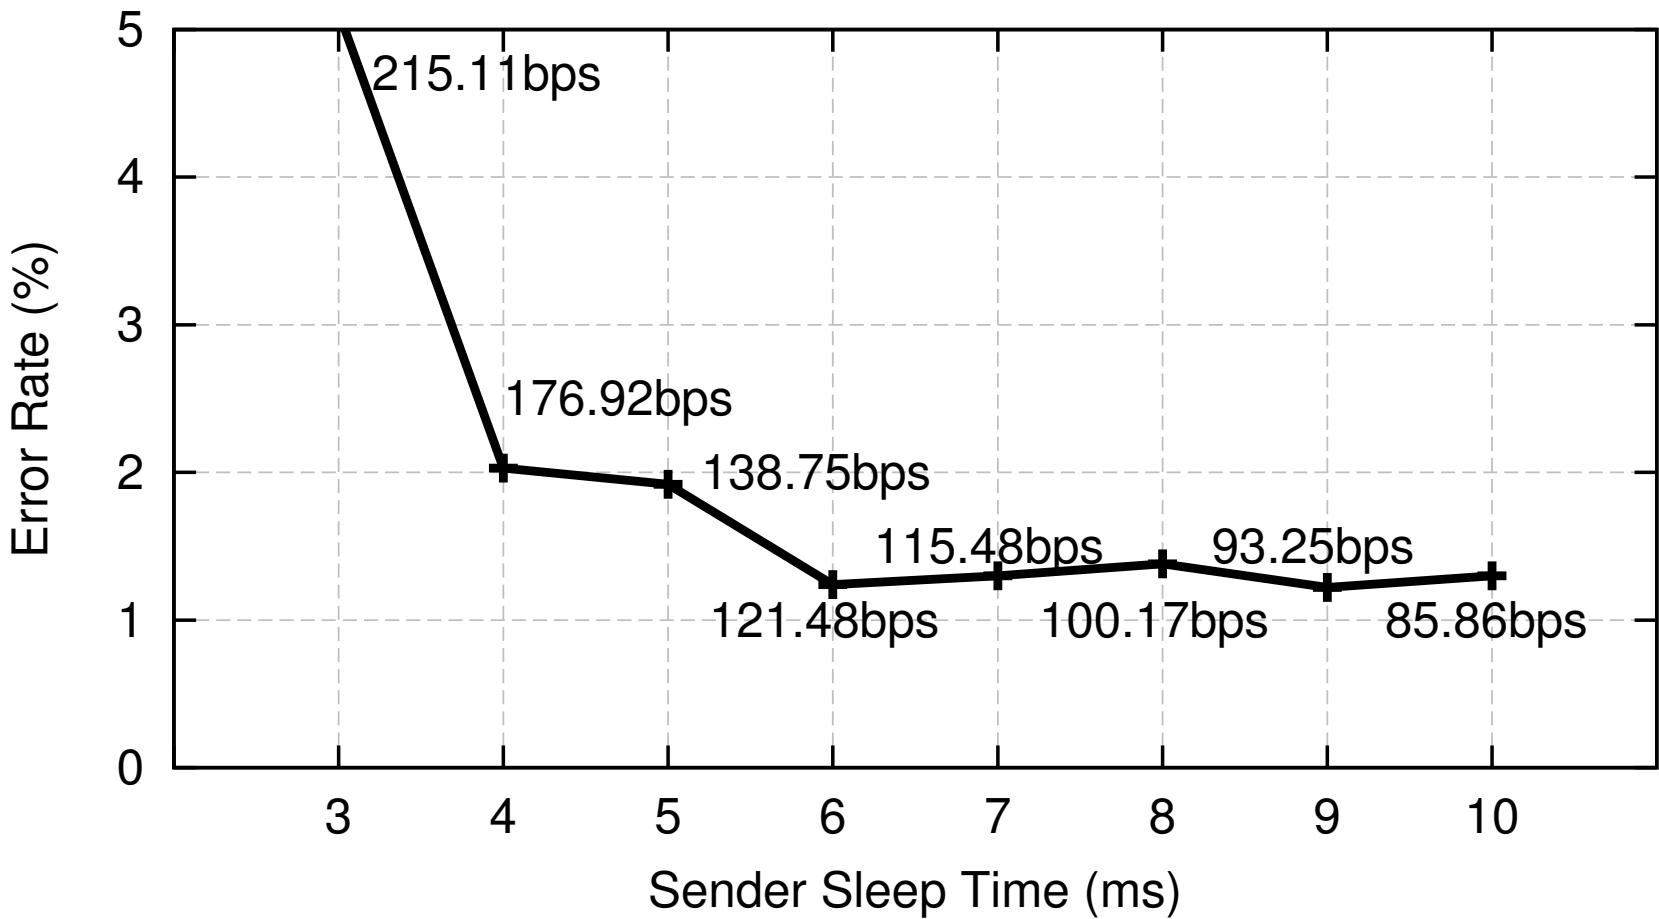
$$1\text{bit}/(1.47 + 3)\text{ms} \approx 223.71\text{bps}$$



Acknowledge:
Busy loops until CPU counter has big jump

Profiled read time:
2.34ms

Sleep time vs. Bit rate



Cost of Synchronization

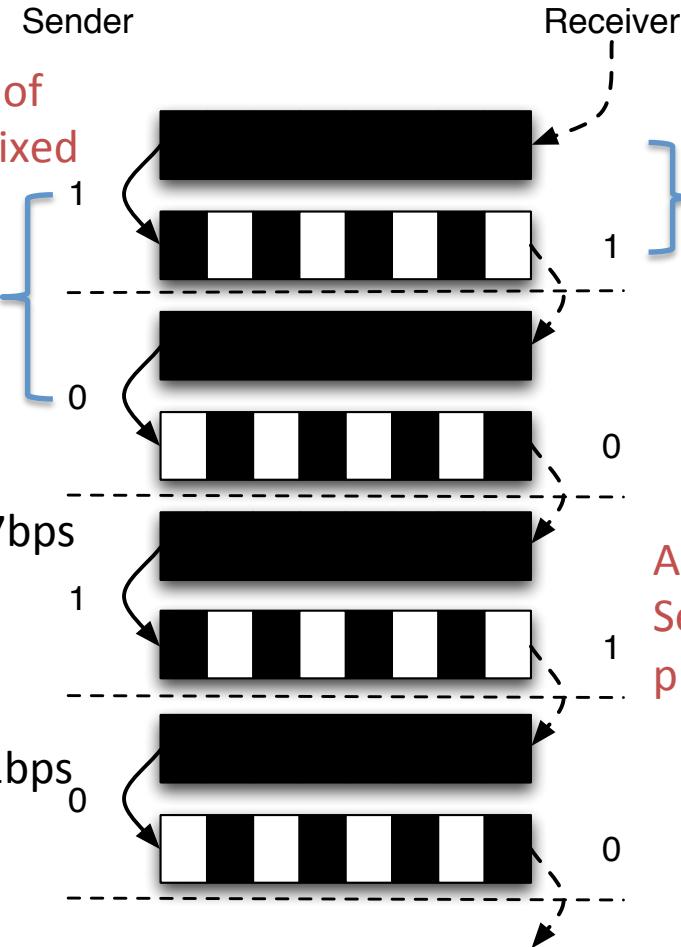
ASSUMPTION: The variation of reading time is small, if not fixed

Timing:
Sleep 3ms

$$1\text{bit}/(1.47 + 2.34)\text{ms} \approx 262.47\text{bps}$$



$$1\text{bit}/(1.47 + 3)\text{ms} \approx 223.71\text{bps}$$



Acknowledge:
Busy loops until CPU
counter has big jump

ASSUMPTION:
Sender and receiver share the
physical CPU

CPU Caps

- 40% hard limit on EC2 small instances
 - Cannot communicate in 60% of time

$$1\text{bit}/(1.47\text{ms} + 3\text{ms})^* \approx 223.71\text{bps}$$



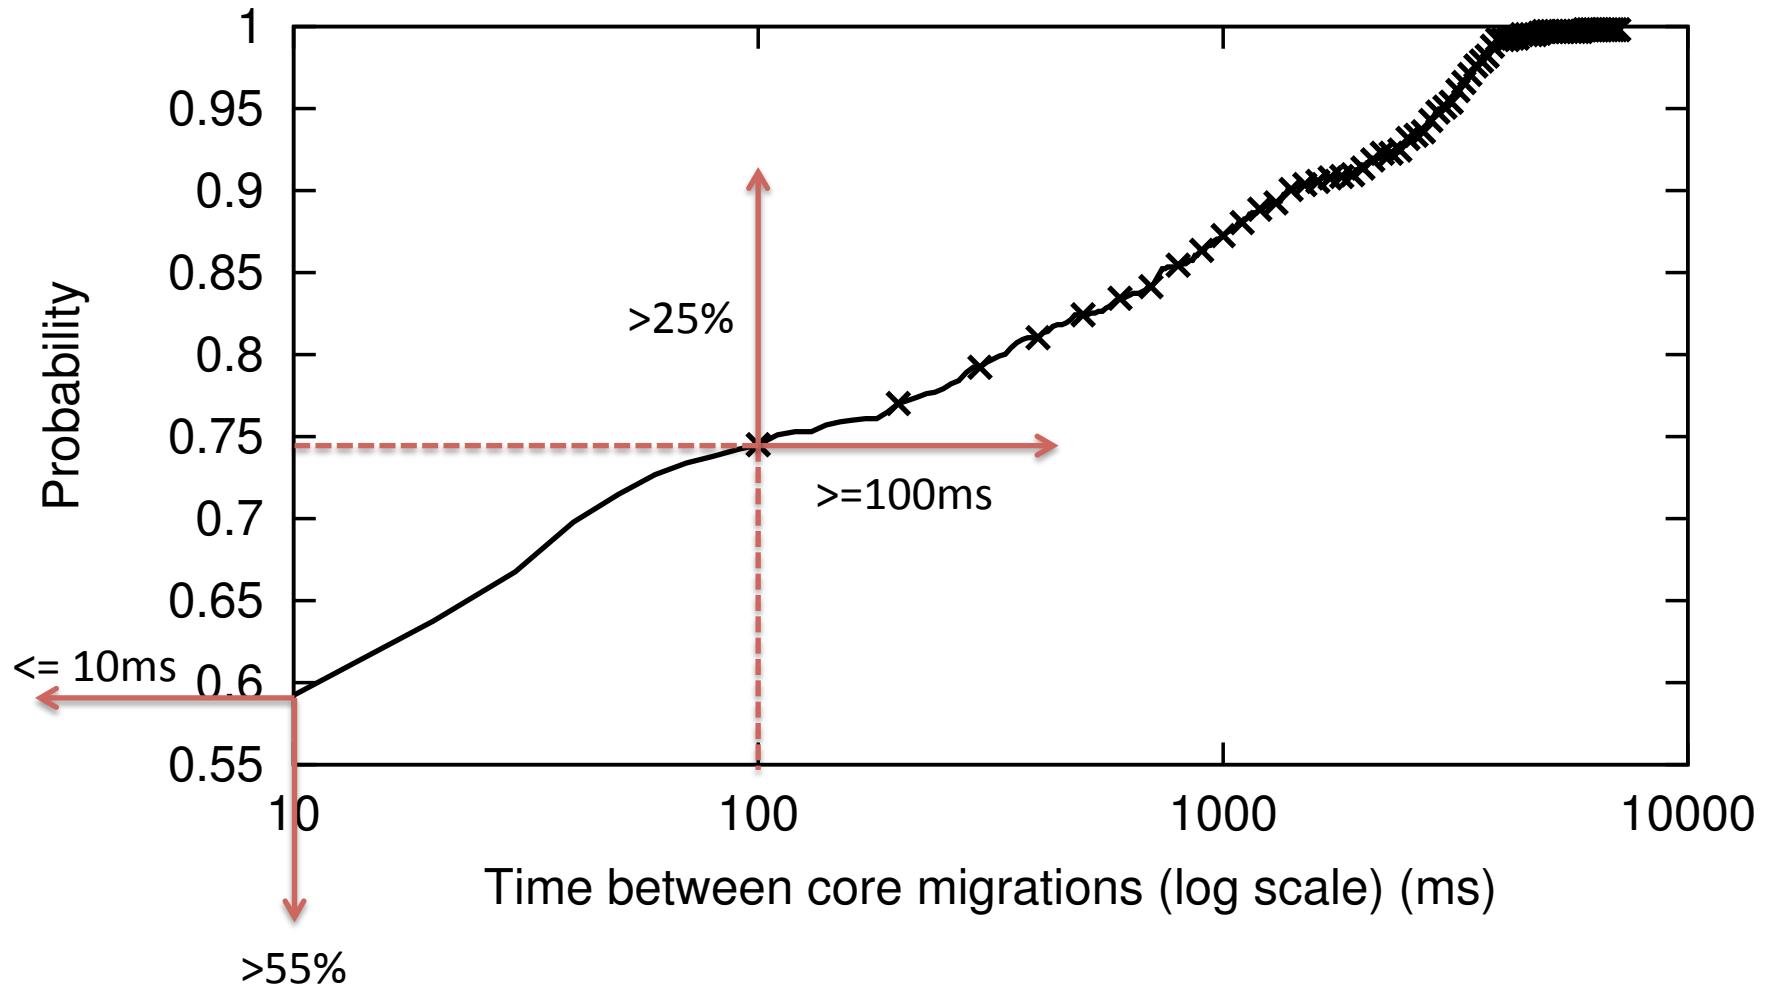
$$1\text{bit}/(1.47\text{ms} + 3\text{ms}) * 40\% \approx 89.48\text{bps}$$

* Refer to the paper for the numbers on EC2

Core Migration

- EC2 VCPUs are not pinned on physical CPU
 - CPU counter trick requires shared CPU
 - Two physical CPUs may not share L2 cache
- EC2's core migration policy is unclear
 - Empirical data

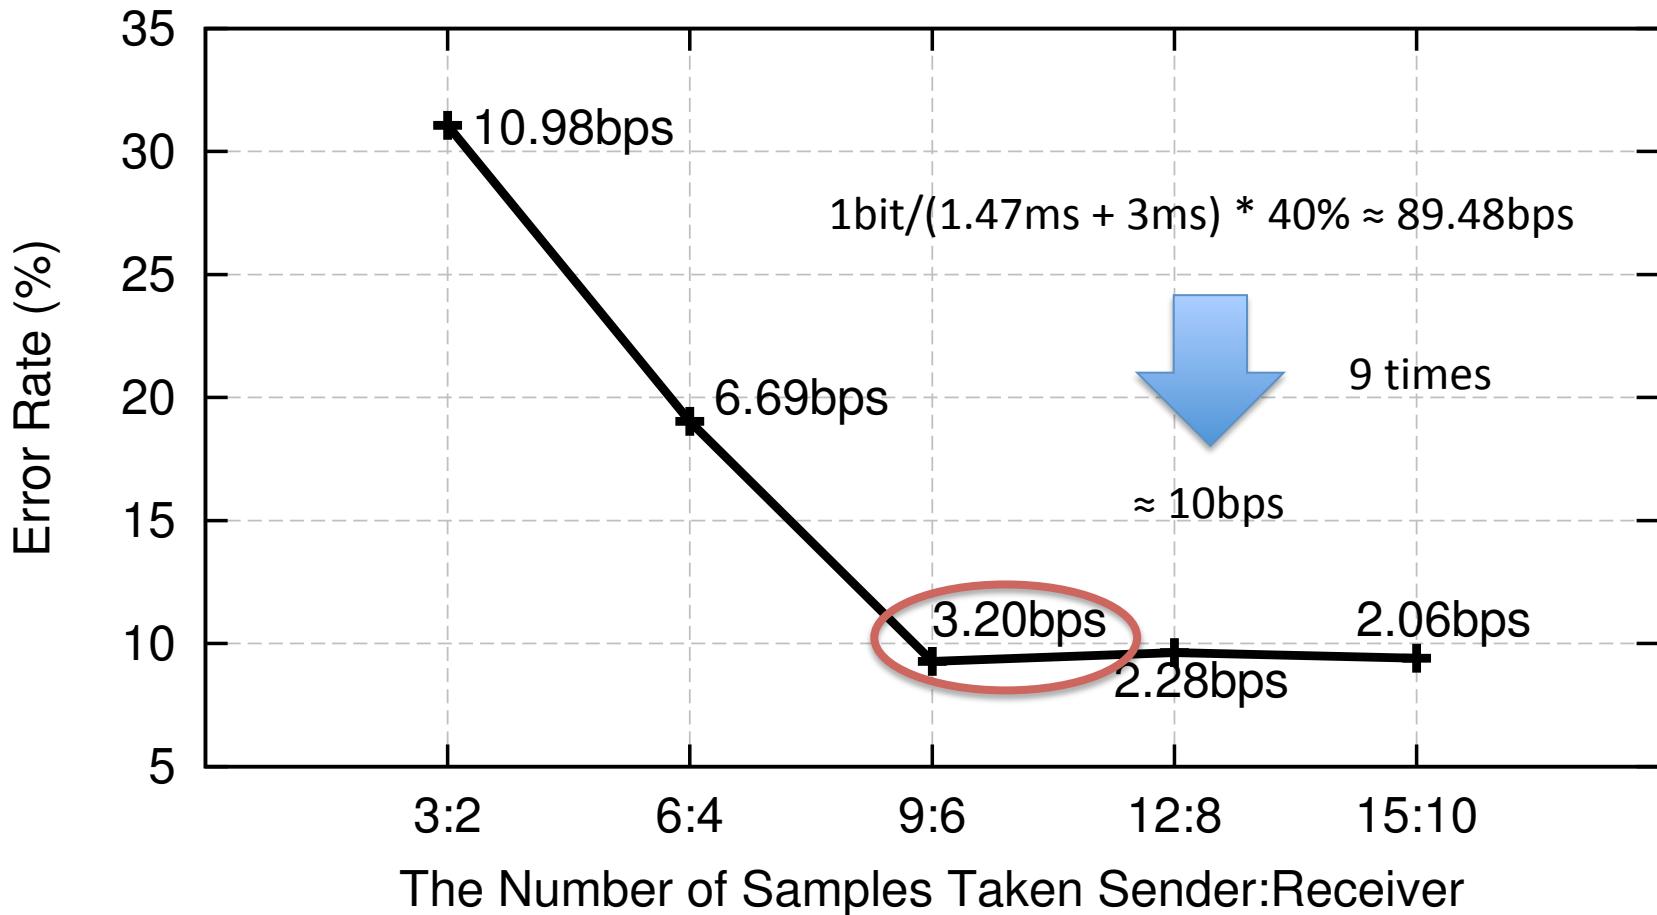
Profile the Core Migration



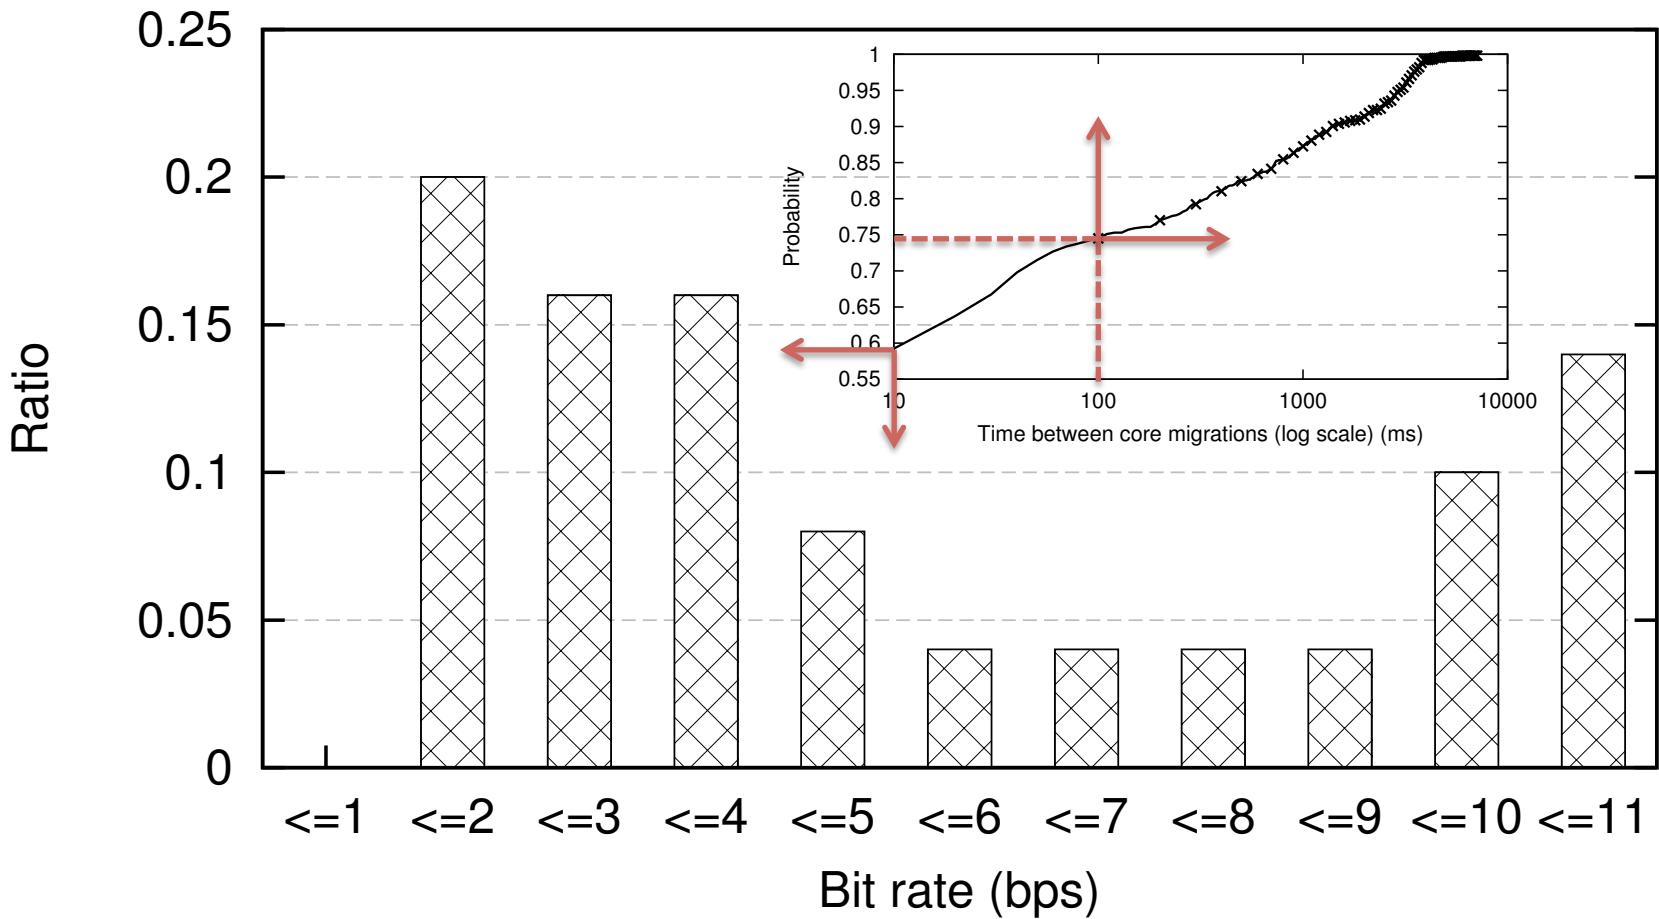
Sharing resistant Synchronization

- Ristenpart et al.
 - Take multiple samples to make a decision
 - Revert to the beginning if core migrated
- Challenges of this approach:
 - Effectiveness of the CPU counter technique
 - Repeat: how many times?
- A protocol is detailed on the paper
 - Use L2 cache itself to indicate activities
 - Explore optimal number of repeats

The Number of Samples Taken



Bit Rate Distribution



Conclusion and Future Work

- Questions to be answered
 - Maximum achievable bit rate
 - Factors influence the bit rate
- A better picture
 - Sleep time little impact if selected right
 - 40% CPU cap: ~200bps -> ~90bps
 - Core migration: ~90bps -> 10bps max, 3bps average
- Data exfiltration using covert channels on EC2
 - No documents leaked out to wikileaks this way
 - Leak small secrets like private keys
- Future work
 - Other cloud providers

Flipping Bits in Memory Without Accessing Them

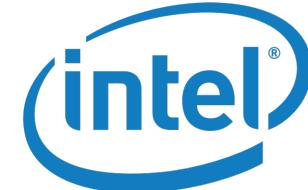
DRAM Disturbance Errors

Yoongu Kim

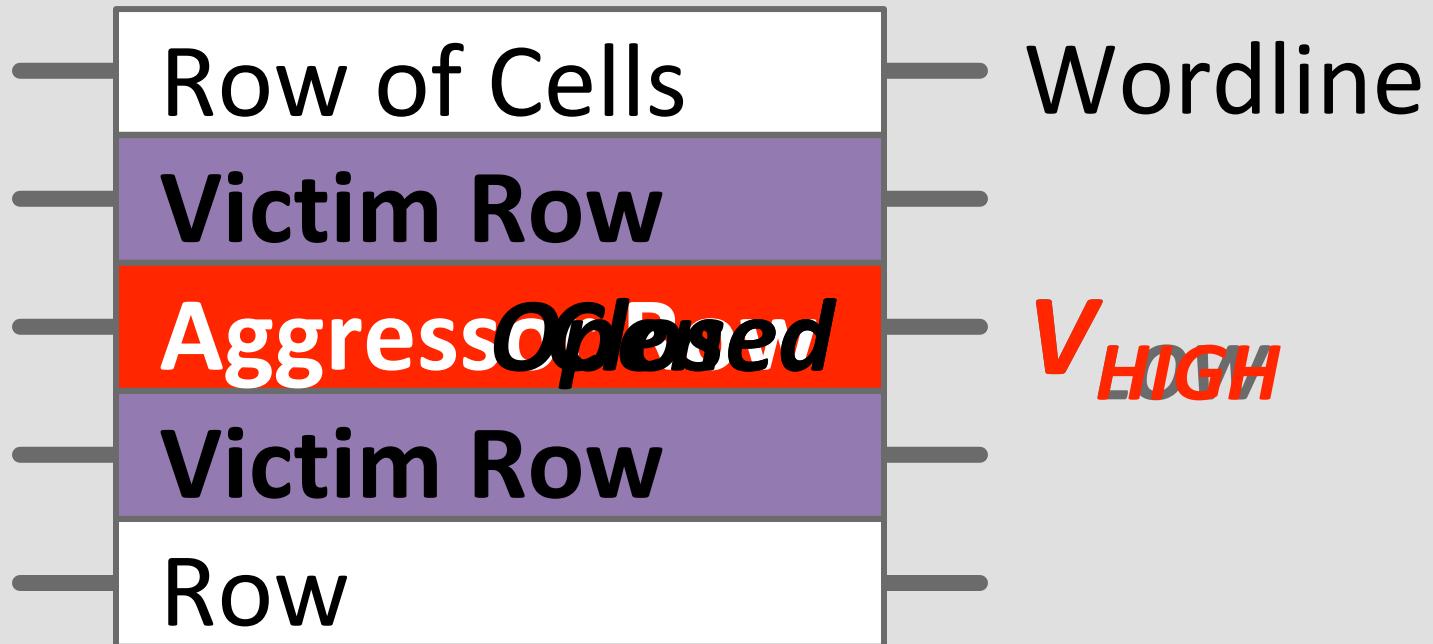
Ross Daly, Jeremie Kim, Chris Fallin, Ji Hye Lee,
Donghyuk Lee, Chris Wilkerson, Konrad Lai, Onur Mutlu

Carnegie Mellon

SAFARI



DRAM Chip



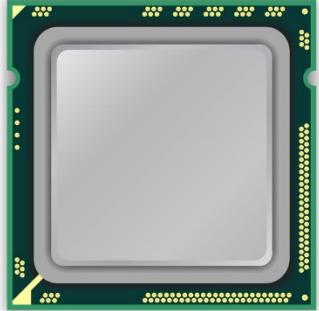
*Repeatedly opening and closing a row induces **disturbance errors** in adjacent rows*

Quick Summary of Paper

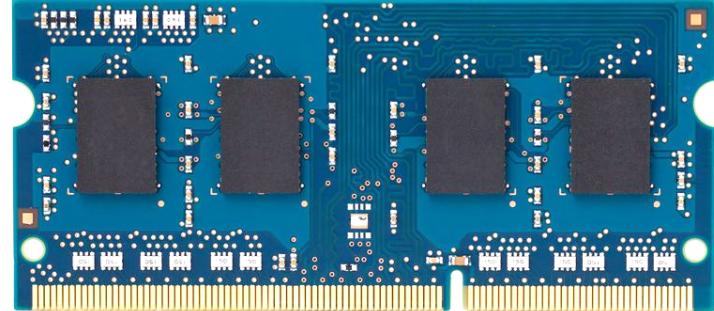
- We expose the **existence and prevalence** of disturbance errors in DRAM chips of today
 - 110 of 129 modules are vulnerable
 - Affects modules of 2010 vintage or later
- We characterize the **cause and symptoms**
 - Toggling a row accelerates charge leakage in adjacent rows: *row-to-row coupling*
- We prevent errors using a **system-level approach**
 - Each time a row is closed, we refresh the charge stored in its adjacent rows with a low probability

How to Induce Errors

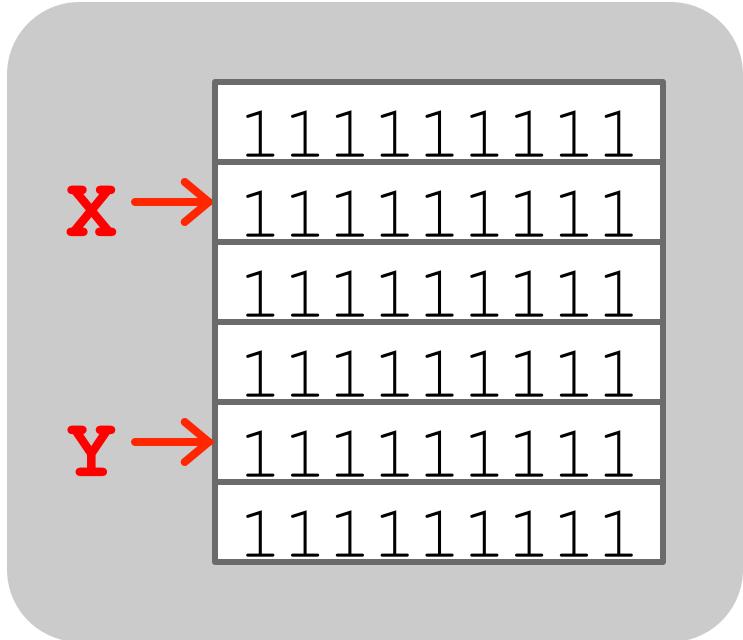
x86 CPU DRAM Module



The logo consists of the text "DDR3" in a bold, white, sans-serif font, centered within a dark gray double-headed arrow graphic.



1. Avoid *cache hits*
 - Flush **X** from cache
 2. Avoid *row hits* to **X**
 - Read **Y** in another row

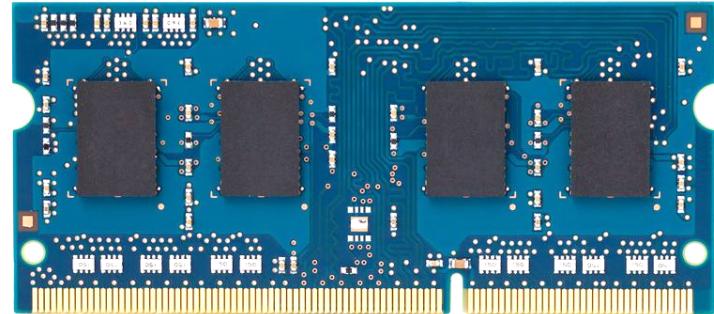


How to Induce Errors

x86 CPU DDRAM Module



DDR3



```
loop:  
    mov (%X), %eax  
    mov (%Y), %ebx  
    clflush (%X)  
    clflush (%Y)  
    mfence  
    jmp loop
```

X →	001110111
	1111 1111
	101111101
	110001011
Y →	1111 1111
	011011110

Security Implications

- *Breach of memory protection*
 - OS page (4KB) fits inside DRAM row (8KB)
 - Adjacent DRAM row → Different OS page
- *Vulnerability: disturbance attack*
 - By accessing its own page, a program could corrupt pages belonging to another program
- *We constructed a proof-of-concept*
 - Using only user-level instructions

Exploiting the DRAM rowhammer bug to gain kernel privileges

How to cause and exploit
single bit errors

Mark Seaborn and Thomas Dullien

Kernel exploit

- x86 page tables entries (PTEs) are **dense and trusted**
 - They control access to physical memory
 - A bit flip in a PTE's physical page number can give a process access to a different physical page
- Aim of exploit: Get access to a page table
 - Gives access to all of physical memory
- Maximise chances that a bit flip is useful:
 - Spray physical memory with page tables
 - Check for useful, repeatable bit flip first

x86-64 Page Table Entries (PTEs)

- Page table is a 4k page containing array of 512 PTEs
- Each PTE is 64 bits, containing:

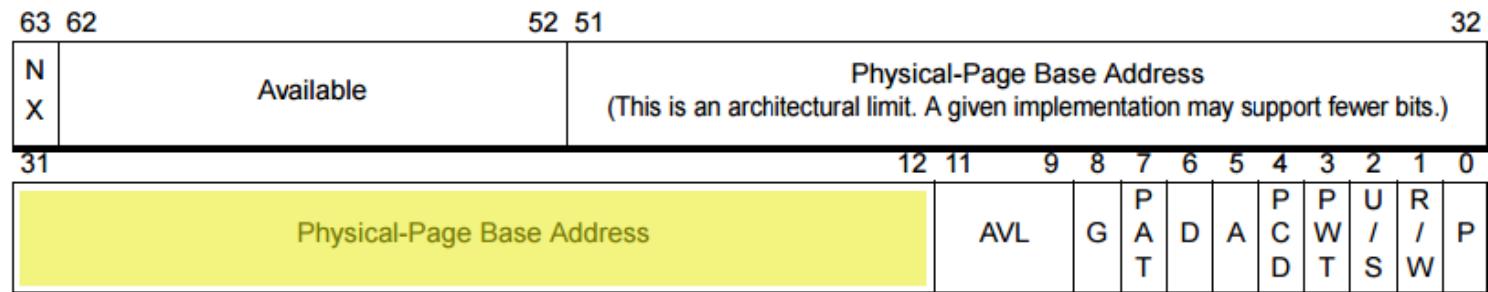
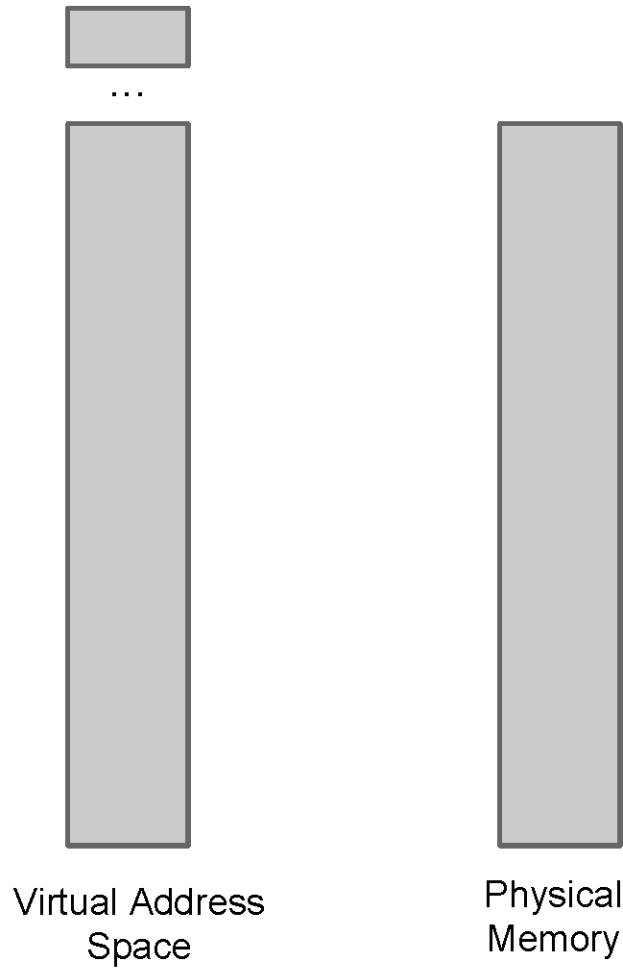
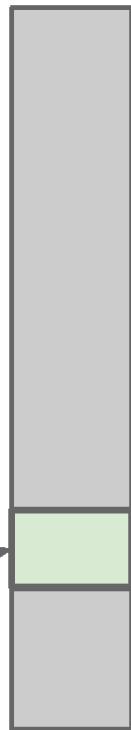
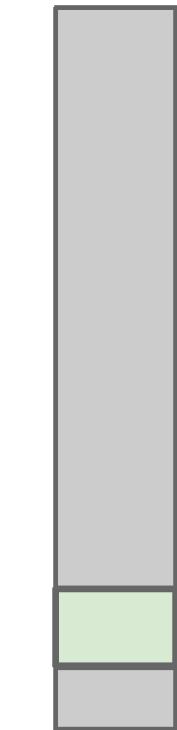


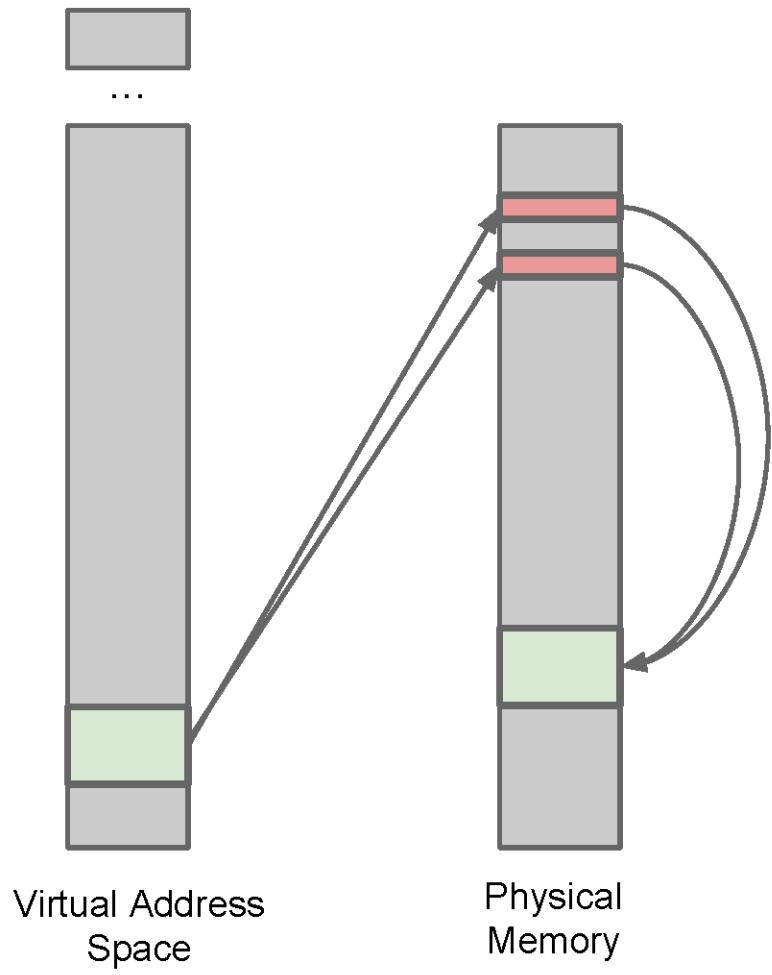
Figure 5-21. 4-Kbyte PTE—Long Mode

- Could flip:
 - “Writable” permission bit (RW): 1 bit → 2% chance
 - Physical page number: 20 bits on 4GB system → 31% chance





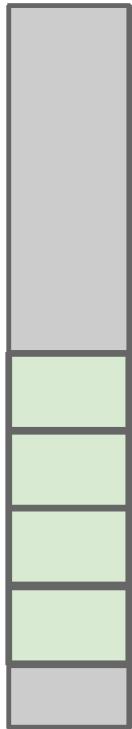
What happens when we map a file with read-write permissions?



What happens when we map a file with read-write permissions? Indirection via page tables.



...

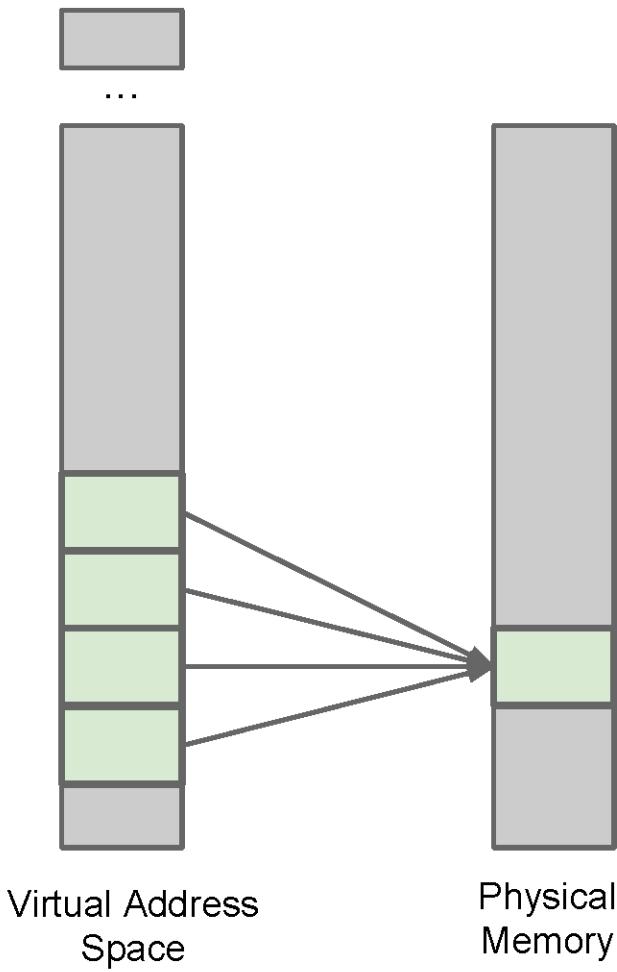


Virtual Address
Space



Physical
Memory

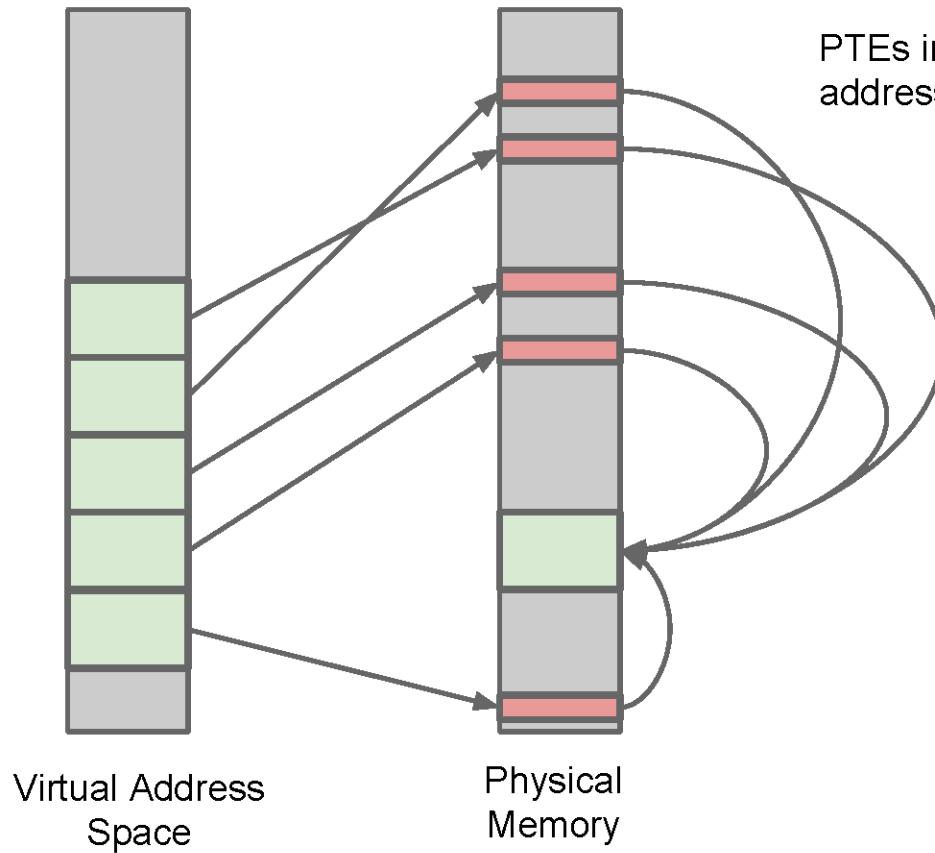
What happens when we repeatedly map a file with
read-write permissions?



What happens when we repeatedly map a file with
read-write permissions?



...

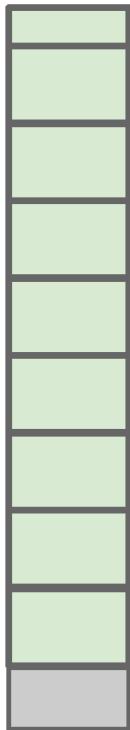


What happens when we repeatedly map a file with read-write permissions?

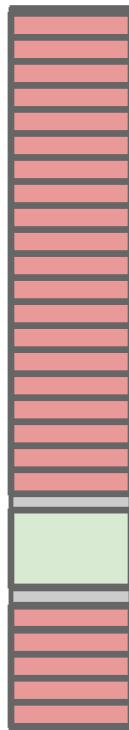
PTEs in physical memory help resolve virtual addresses to physical pages.



...



Virtual Address
Space

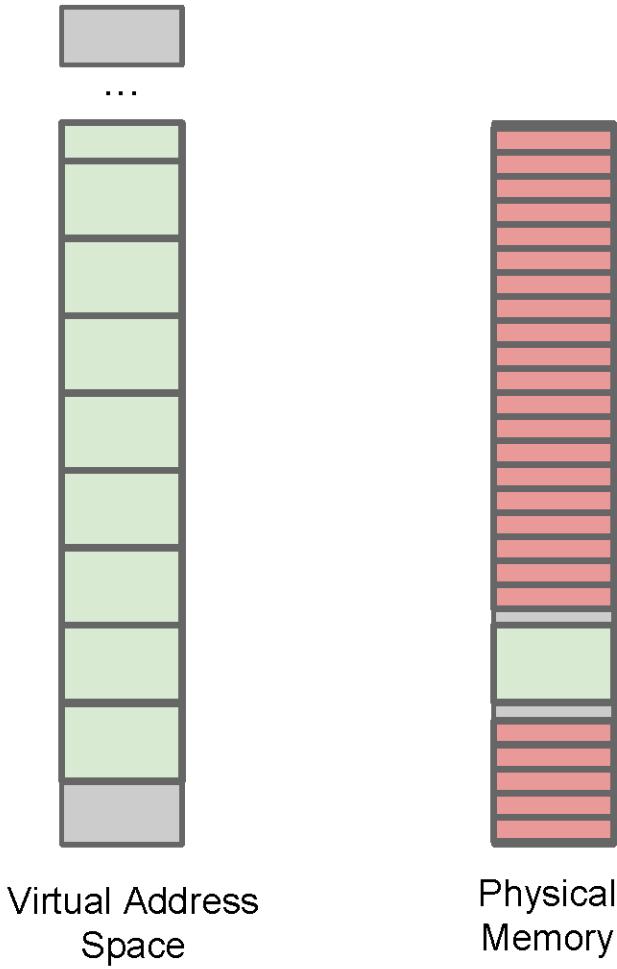


Physical
Memory

What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

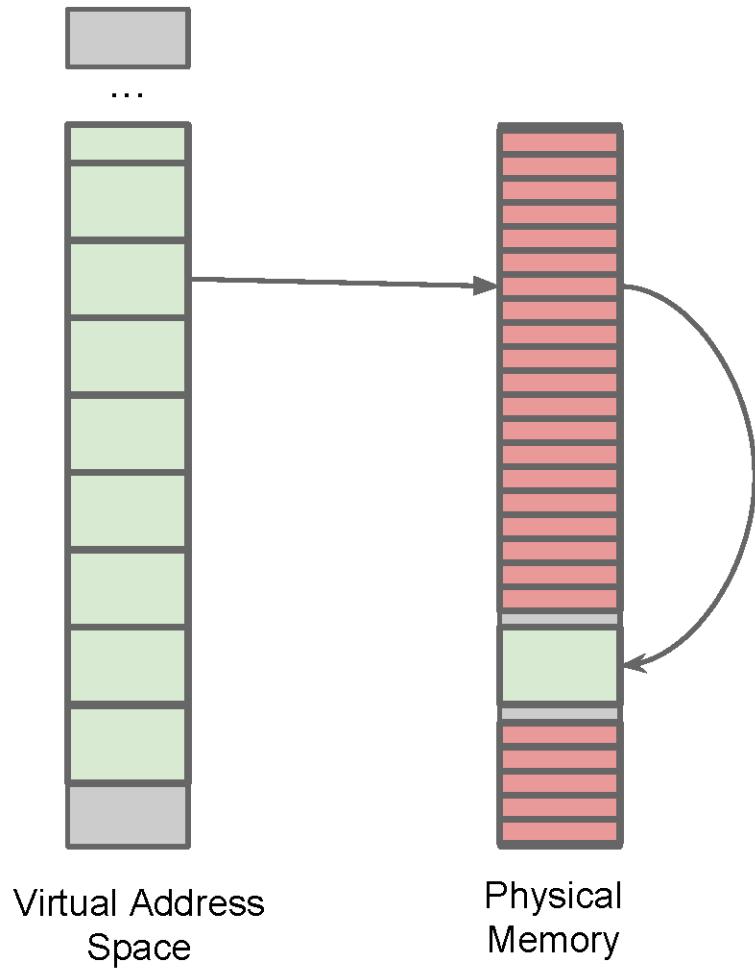


What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

Each of them points to pages in the same physical file mapping.



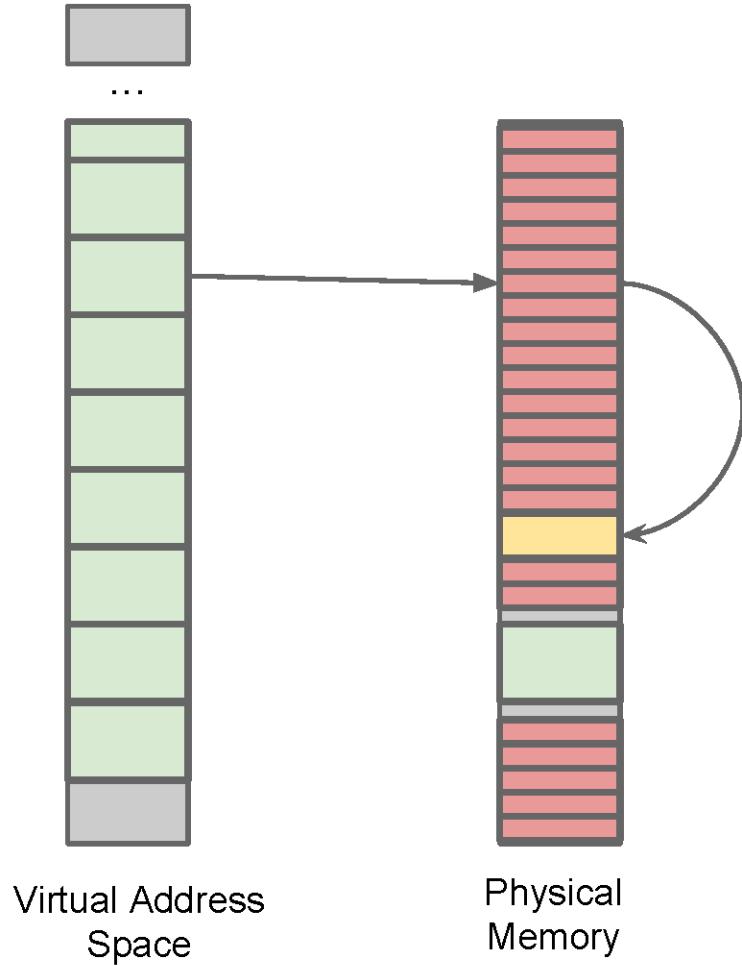
What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

Each of them points to pages in the same physical file mapping.

If a bit in the right place in the PTE flips ...



What happens when we repeatedly map a file with read-write permissions?

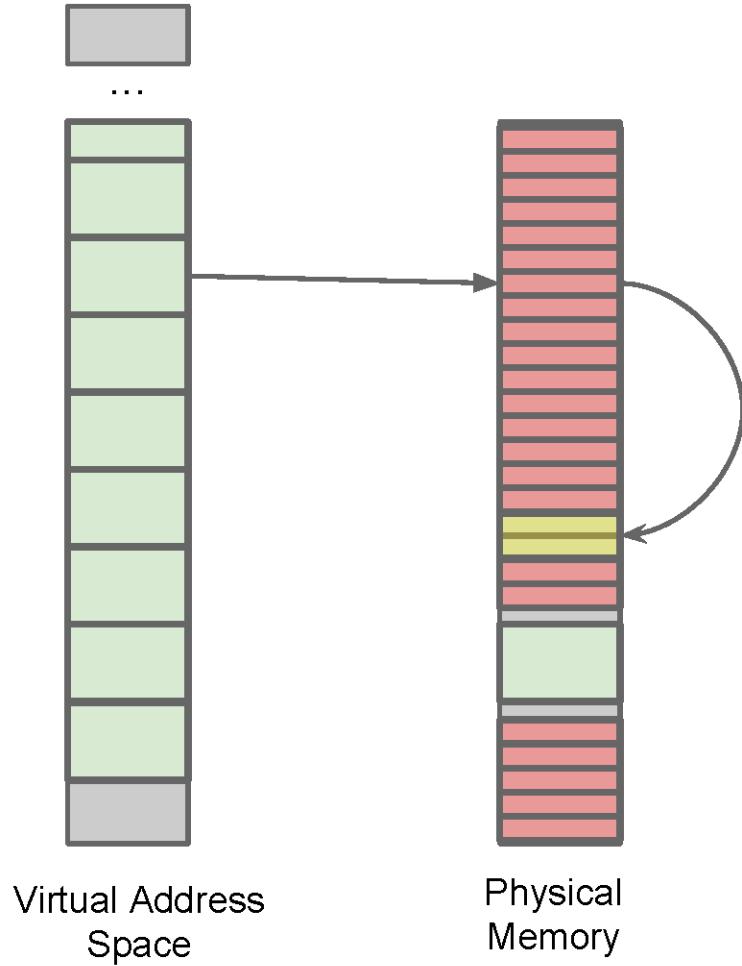
PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

Each of them points to pages in the same physical file mapping.

If a bit in the right place in the PTE flips ...

... the corresponding virtual address now points to a wrong physical page - with RW access.



What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

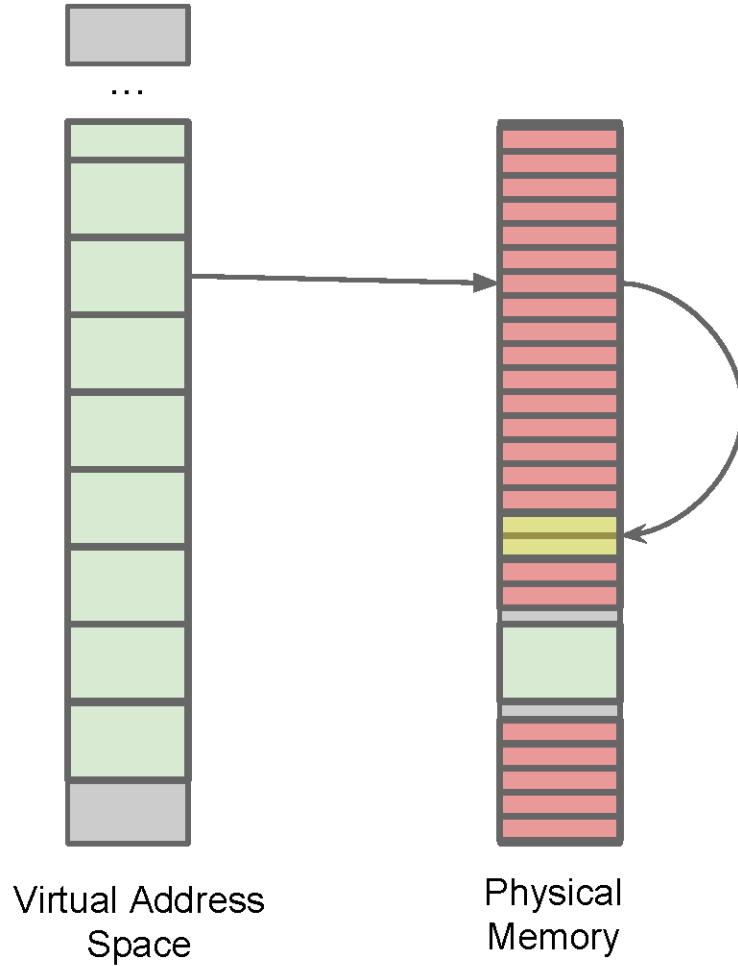
We can fill physical memory with PTEs.

Each of them points to pages in the same physical file mapping.

If a bit in the right place in the PTE flips ...

... the corresponding virtual address now points to a wrong physical page - with RW access.

Chances are this wrong page contains a page table itself.



What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

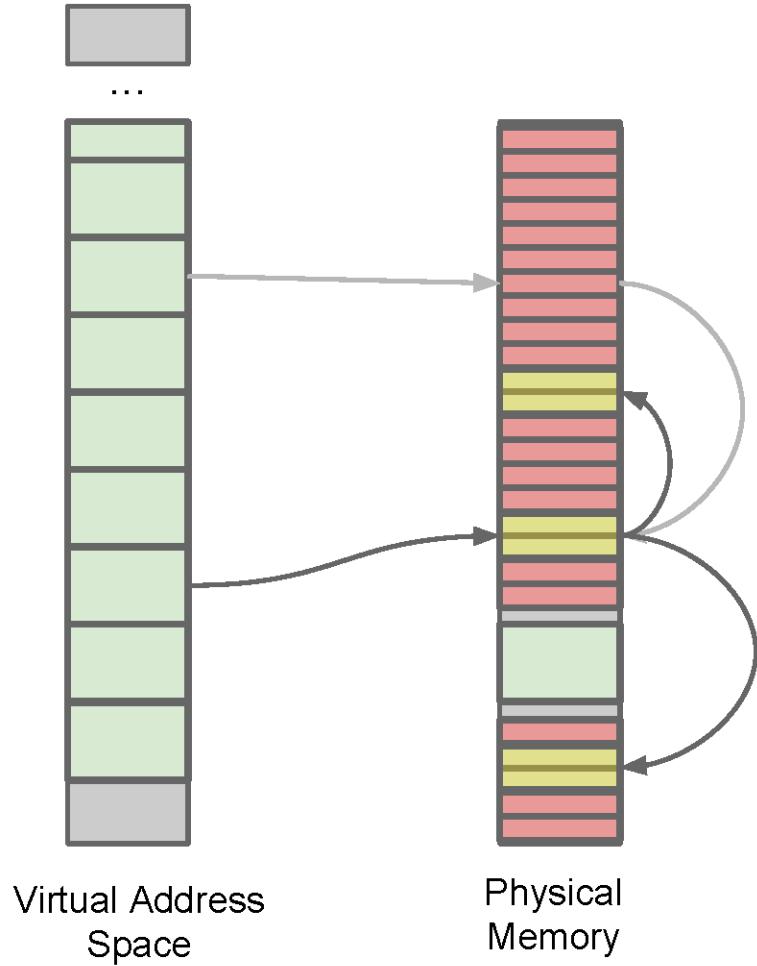
Each of them points to pages in the same physical file mapping.

If a bit in the right place in the PTE flips ...

... the corresponding virtual address now points to a wrong physical page - with RW access.

Chances are this wrong page contains a page table itself.

An attacker that can read / write page tables ...



What happens when we repeatedly map a file with read-write permissions?

PTEs in physical memory help resolve virtual addresses to physical pages.

We can fill physical memory with PTEs.

Each of them points to pages in the same physical file mapping.

If a bit in the right place in the PTE flips ...

... the corresponding virtual address now points to a wrong physical page - with RW access.

Chances are this wrong page contains a page table itself.

An attacker that can read / write page tables can use that to map **any** memory read-write.

Exploit strategy

Privilege escalation in 7 easy steps ...

1. Allocate a large chunk of memory
2. Search for locations prone to flipping
3. Check if they fall into the “right spot” in a PTE for allowing the exploit
4. Return that particular area of memory to the operating system
5. Force OS to re-use the memory for PTEs by allocating massive quantities of address space
6. Cause the bitflip - shift PTE to point into page table
7. Abuse R/W access to all of physical memory

In practice, there are many complications.